

# Clustering in Computer Vision

12 October 2018

CSE473/573 – Autumn 2018


# Clustering

- What is clustering?
  - Grouping of “objects” into meaningful categories
  - Given a representation of  $N$  objects, find  $k$  clusters based on a suitable measure of similarity.
- Data Clustering is useful in and beyond Computer Vision
  - Segmentation as clustering (today)
  - Texture modeling
  - Quantization
  - Beyond
    - Data exploration
    - Compression
    - Natural classification

# Feature Space


- Every token is identified by a set of salient visual characteristics. For example:
  - Position
  - Color
  - Texture
  - Motion vector
  - Size, orientation (if token is larger than a pixel)

# Feature Space - Example




# Feature Space - Example

*Feature space:*  
each token is represented by a point



# Feature Space - Similarity

Token **similarity** is thus measured by distance between points (“feature vectors”) in feature space




$$\sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \cdots + (p_n - q_n)^2} = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$


Source: Savarese slides.

# Feature Space - Similarity

Cluster together tokens with high similarity




# Color based Clustering - Example




©: Savarese slides.



# Color based Clustering - Example




# Color based Clustering - Example



source: Savarese slides.

# Color based Clustering - Example



Source: Savarese slides.

# Clustering – formal definition

- Given a set of  $N$  data samples  $D = x_1, x_2, \dots, x_N$  in a  $d$ -dimensional feature space,  $D$  is partitioned into a number of disjoint subsets  $D_j$ :

$$D = \bigcup_{j=1}^k D_j \quad \text{where} \quad D_i \cap D_j = \emptyset \quad \forall i \neq j$$

where the points in each subset are similar to each other according to the given similarity function.

- A partition is denoted by

$$\pi = (D_1, D_2, \dots, D_k)$$

and clustering is then formulated as


$$\pi^* = \arg \min_{\pi} f(\pi)$$

for  $f(\cdot)$  that captures the desired cluster properties.

# K-Means Clustering

Iterative process. Takes multiple iterations to converge


1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$



# Initialize K centers

## K-Means Clustering


1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$



# Assign nearest points to clusters

## K-Means Clustering

1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$






# Re-compute centers of clusters

## K-Means Clustering

1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$




Recompute centers



# Re-Assign points to clusters

1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$




Assign points to closest centers

# Iterate....

## K-Means Clustering

1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$




Assign points to closest centers

# Iterate....

## K-Means Clustering

1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$




Recompute centers

# Stop! When the centers do not move

## K-Means Clustering

1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$




Points already assigned to nearest  
centers: Algorithm ends

# Stop! When the centers do not move

## K-Means Clustering

1. Randomly initialize  $\mu_1, \mu_2, \dots, \mu_c$
2. Repeat until no change in  $\mu_i$ :
  - (a) Classify  $N$  samples according to nearest  $\mu_i$
  - (b) Recompute  $\mu_i$



Points already assigned to nearest  
centers: Algorithm ends

# K-Means Clustering Animation

<http://shabal.in/visuals/kmeans/1.html>

<http://simplystatistics.org/2014/02/18/k-means-clustering-in-a-gif/>

# K-Means Clustering – Pros and Cons


- Pros

- Simple and fast
- (Always) converges to a local minimum of the error function
- Available implementations (e.g., in Matlab)



$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

- Cons

- Need to pick K
- Sensitive to initialization
- Only finds “spherical” clusters
- Sensitive to outliers




# K-Means Clustering – Example






# What are the clusters here?



# K-Means Clustering fails




# Challenges


1. What is a cluster?
2. How to define pair-wise similarity?
3. Which features? Which normalizations scheme?
4. How many clusters?
5. Which clustering method?
6. Are the discovered clusters and partitioning valid?
7. Does the data have any clustering tendency?

# Cluster Validity

- Clustering algorithms find clusters, even if there are no **natural** clusters in the data.



100 2D uniform data points



k-Means with  $k=3$

# Soft vs Hard Clustering

- Kmeans performs Hard clustering:
  - Data point is deterministically assigned to one and only one cluster
  - But in reality clusters may overlap
- Soft-clustering:
  - Data points are assigned to clusters with certain probabilities

# Probabilistic Clustering

- Basic questions
  - what's the probability that a point  $\mathbf{x}$  is in cluster  $m$ ?
  - what's the shape of each cluster?
- K-means doesn't answer these questions
- Basic idea
  - instead of treating the data as a bunch of points, assume that they are all generated by sampling a continuous function
  - This function is called a **generative model**

# How can we extend K-means to make soft clustering


- Given a set of clusters centers  $\mu_1, \mu_2, \dots, \mu_k$ , instead of directly assign all data points to their closest clusters, we can assign them **partially (probabilistically) based on the distances**
- This can be done by
  - assuming a probabilistic distribution (model) for each cluster
  - compute the probability that each point belongs to each cluster
  - Often referred to as Model-based clustering

# Gaussian for representing a cluster

- What exactly is a cluster?
  - Intuitively it is a tightly packed ball-shape like thing
- We can use a Gaussian (normal) distribution to describe it
- Let's first review what is a Gaussian distribution




# Gaussian Distribution



Probability density function  $f(x)$  is a function of  $x$  given  $\mu$  and  $\sigma$

$$N(x | \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2\right)$$

# Gaussian Distribution




Define likelihood as a function of  $\mu$  and  $\sigma$   
given  $x_1, x_2, \dots, x_n$

$$\prod_{i=1}^n N(x_i | \mu, \sigma^2)$$

# Multi-variate Gaussian

- Univariate Gaussian distribution:





$$N(\mu, \sigma^2)$$

$\mu$  – mean, center of the mass

$\sigma^2$  – standard deviation, spread of the mass

- Multivariate Gaussian distribution:




$$N(\mu, \Sigma)$$

$\mu$  –  $(\mu_1, \mu_2)$

$\Sigma$  – Covariance matrix


$$\begin{bmatrix} \sigma_1^2 & \sigma_{12} \\ \sigma_{12} & \sigma_2^2 \end{bmatrix}$$

# Different Co-Variance Matrices




**Covariance matrix  $\Sigma =$**

$$\begin{matrix} \sigma^2 & 0 \\ 0 & \sigma^2 \end{matrix}$$



**Covariance matrix  $\Sigma =$**

$$\begin{matrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{matrix}$$



**Covariance matrix  $\Sigma =$**


$$\begin{matrix} \sigma_1^2 & \sigma_{12} \\ \sigma_{12} & \sigma_2^2 \end{matrix}$$

# Log-Likelihood

- Multivariate Gaussian

$$\mathcal{N}(x|\mu, \Sigma) = \frac{1}{(2\pi|\Sigma|)^{1/2}} \exp \left\{ -\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu) \right\}$$

mean                      covariance



- Log likelihood

$$L(\mu, \Sigma) = \sum_{i=1}^n \ln N(x_i | \mu, \Sigma) = \sum_{i=1}^n \left( -\frac{1}{2}(x_i - \mu)^T \Sigma^{-1}(x_i - \mu) \right) - \pi \ln |\Sigma|$$

# Gaussian Mixture Models (GMMs)

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2}|\boldsymbol{\Sigma}|^{1/2}} \exp \left[ -\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}) \right]$$





- It forms the basis for the mixture of Gaussians density
- The Gaussian mixture is linear superposition of Gaussians:

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- The  $\pi_k$  are non-negative scalars called mixing coefficients and they govern the relative importance between the various Gaussians in the mixture density.  $\sum_k \pi_k = 1$

# Gaussian Mixture Models (GMMs)

In below examples – the distribution is a mixture of three gaussians with different means and variances



# Clustering using GMMs

- Given a set of data points, and assume that we know there are  $k$  clusters in the data, we need to:
  - Assign the data points to the  $k$  clusters (soft assignment)
  - Learn the gaussian distribution parameters for each cluster:  $\mu$  and  $\Sigma$




# Expectation Maximization (EM)

- Randomly initialize the Gaussian parameters
- Repeat until converge
  1. Compute  $P(\mathbf{x}^j \in C_i | \mathbf{x}^j)$  for all data points and all clusters


This is called the E-step for it computes the expected values of the cluster memberships for each data point
  2. Re-compute the parameters of each Gaussian

This is called the M-step for it performs maximum likelihood estimation of parameters


# GMM Example - Initialization




# After First Iteration




# After Second Iteration



## After Sixth Iteration



## After 20<sup>th</sup> Iteration



Q: Why are these two points red when they appear to be closer to blue?


# K-Means vs GMMs

- Objective function
    - Minimize sum of squared Euclidean distance
  - Can be optimized by an EM algorithm
    - E-step: assign points to clusters
    - M-step: optimize clusters
    - Performs hard assignment during E-step
  - Assumes spherical clusters with equal probability of a cluster
- Objective function
    - Maximize log-likelihood
  - EM algorithm
    - E-step: Compute posterior probability of membership
    - M-step: Optimize parameters
    - Perform soft assignment during E-step
  - Can be used for non-spherical clusters
  - Can generate clusters with different probabilities

# Spectral Clustering

Spectral clustering techniques make use of the spectrum (eigenvalues) of the similarity matrix of the data to perform dimensionality reduction before clustering in fewer dimensions

- **Spectral approach**
  - Use similarity graphs to encode local neighborhood information
  - Data points are vertices of the graph
  - Connect points which are “close”






# Similarity Graph

- Represent dataset as a weighted graph  $G(V,E)$
- All vertices which can be reached from each other by a path form a connected component
- Only one connected component in the graph—The graph is fully connected

$V=\{x_i\}$  Set of  $n$  vertices representing data points

$E=\{W_{ij}\}$  Set of weighted edges indicating pair-wise similarity between points




# Clustering Objective

## Traditional definition of a “good” clustering

- Points assigned to same cluster should be highly similar
- Points assigned to different clusters should be highly dissimilar


## Apply this objective to our graph representation



Minimize weight of between-group connections


# Graph Cut

- Express clustering objective as a function of the **edge cut** of the partition
- Cut**: Sum of weights of edges with only one vertex in each group
- We want to find the **minimal cut** between groups



# Minimum Cut - Example

- Minimum Cut



$$\text{Cut}(BCDE, A) = 0.17$$

# Problem


- Identifying a minimum cut is NP-hard
- There are efficient approximations using linear algebra
- Based on the Laplacian Matrix, or **graph Laplacian**

# Similarity Matrix Representation

- **Similarity matrix ( $W$ )**

- $n \times n$  matrix

- $W = [w_{ij}]$  : edge weight between vertex  $x_i$  and  $x_j$



	$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$x_6$
$x_1$	0	0.8	0.6	0	0.1	0
$x_2$	0.8	0	0.8	0	0	0
$x_3$	0.6	0.8	0	0.2	0	0
$x_4$	0	0	0.2	0	0.8	0.7
$x_5$	0.1	0	0	0.8	0	0.8
$x_6$	0	0	0	0.7	0.8	0

- **Important properties**


- Symmetric matrix

# Degree Matrix

- **Degree matrix ( $D$ )**

- $n \times n$  diagonal matrix

- $D(i,i) = \sum_j w_{ij}$  : total weight of edges incident to vertex  $x_i$



	$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$x_6$
$x_1$	1.5	0	0	0	0	0
$x_2$	0	1.6	0	0	0	0
$x_3$	0	0	1.6	0	0	0
$x_4$	0	0	0	1.7	0	0
$x_5$	0	0	0	0	1.7	0
$x_6$	0	0	0	0	0	1.5

- **Used to**


- Normalize adjacency matrix

# Laplacian Matrix

- Laplacian matrix ( $L$ )**

—  $n \times n$  symmetric matrix

$$L = D - W$$



	$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$x_6$
$x_1$	1.5	-0.8	-0.6	0	-0.1	0
$x_2$	-0.8	1.6	-0.8	0	0	0
$x_3$	-0.6	-0.8	1.6	-0.2	0	0
$x_4$	0	0	-0.2	1.7	-0.8	-0.7
$x_5$	-0.1	0	0	-0.8	1.7	-0.8
$x_6$	0	0	0	-0.7	-0.8	1.5



# Find an Optimal mincut

- Express a bi-partition  $(C_1, C_2)$  as a vector

$$f_i = \begin{cases} 1 & \text{if } x_i \in C_1 \\ -1 & \text{if } x_i \in C_2 \end{cases}$$

- We can minimize the cut of the partition by finding a non-trivial vector  $f$  that minimizes the function

$$g(f) = \sum_{i,j \in V} w_{ij} (f_i - f_j)^2 = f^T L f$$

Laplacian  
matrix




# How it works

- How eigen decomposition of  $L$  relates to clustering?

$$\begin{aligned} L &= D - W & f(x_j) &= f_j \text{ cluster assignment} \\ f^T L f &= f^T D f - f^T W f \\ &= \frac{1}{2} \left( \sum_i \left( \sum_j w_{ij} \right) f_i^2 - 2 \sum_{ij} f_i f_j w_{ij} + \sum_j \left( \sum_i w_{ij} \right) f_j^2 \right) \\ &= \frac{1}{2} \sum_{ij} w_{ij} (f_i - f_j)^2 \quad \text{--Cluster objective function} \end{aligned}$$

# Hierarchical Clustering

- Agglomerative approach**



Initialization:

Each object is a cluster


Iteration:

Merge two clusters which are  
most similar to each other;

Until all objects are merged  
into a single cluster

# Hierarchical Clustering

- Divisive Approaches**



Initialization:

All objects stay in one cluster


Iteration:

Select a cluster and split it into  
two sub clusters

Until each leaf cluster contains  
only one object

← Step 4 Step 3 Step 2 Step 1 Step 0 Top-down

# Inter-Cluster Distance



- MIN
- MAX
- Group Average
- Distance Between Centroids
- .....

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						


• Distance Matrix

•

# MIN or Single Link


- **Inter-cluster distance**

- The distance between two clusters is represented by the distance of the closest pair of data objects belonging to different clusters.
- Determined by one pair of points, i.e., by one link in the proximity graph




$$d_{\min}(C_i, C_j) = \min_{p \in C_i, q \in C_j} d(p, q)$$

# MIN




Nested Clusters



Dendrogram

# MAX or Complete Link


- **Inter-cluster distance**
  - The distance between two clusters is represented by the distance of the farthest pair of data objects belonging to different clusters




$$d_{\min}(C_i, C_j) = \max_{p \in C_i, q \in C_j} d(p, q)$$



# MAX or Complete Link




Nested Clusters




Dendrogram

# Limitations of MIN



Original Points




Two Clusters

- Sensitive to noise and outliers


# Limitations of MAX

- Tends to break large clusters



# Limitations of MAX


- Tends to break large clusters



# More Techniques with increased complexity

- **Inter-cluster distance**

- The distance between two clusters is represented by the average distance of all pairs of data objects belonging to different clusters
- Determined by all pairs of points in the two clusters



$$d_{\min}(C_i, C_j) = \text{avg}_{p \in C_i, q \in C_j} d(p, q)$$