



Image Segmentation

CSE473/573

Segmentation of Coherent Regions





Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>



Figure-Ground Segmentation

- Separate the foreground object (figure) from the background (ground)



D. Tsai, M. Flagg, and J. M. Rehg. "Motion Coherent Tracking with Multi-label MRF optimization." BMVC 2010.

Grouping of Similar Neighbors



X. Ren and J. Malik. [Learning a classification model for segmentation](#). ICCV 2003.

Slide by Svetlana Lazebnik

Extension beyond Single Images





J. Strom, A. Richardson, E. Olson. "Graph-based Segmentation for Colored 3D Laser Point Clouds." IROS 2010.



M. Grundmann, V. Kwatra, M. Han, I. Essa. "Efficient Hierarchical Graph-Based Video Segmentation." CVPR 2010.

Toy Example



- These intensities define the three groups.
- We could label every pixel in the image according to which of these primary intensities it is.
 - i.e., *segment* the image based on the intensity feature.
- What if the image isn't quite so simple?

Image Histograms




How many “orange” pixels are in this image?


- This type of question answered by looking at the *histogram*
- A histogram counts the number of occurrences of each color
 - Given an image $F[x, y] \rightarrow RGB$
 - The histogram is $H_F[c] = |\{(x, y) \mid F[x, y] = c\}|$
i.e., for each color value c (x-axis), count # of pixels with that color (y-axis)

Histograms of Grayscale intensities

Image



Histogram



Intensity bins

How Many Modes Are There?



- Easy to see, hard to compute

Histograms of Grayscale intensities

Idea: Break the image into K regions (segments) by

- reducing the number of colors to K and
- assigning each pixel to the closest color

Here's what our image looks like if we use two colors (intensities)



Segmentation as Clustering - Issues

Is broccoli same as a pepper?

- K-means clustering based on intensity or color is essentially vector quantization of the image attributes





Slide by Svetlana Lazebnik

Segmentation as Clustering - Improvement

Depending on what we choose as the *feature space*, we can group pixels in different ways.


Grouping pixels based on **intensity+position** similarity



Both regions are black, but if we also include **position (x,y)**, then we could group the two into distinct segments; way to encode both similarity & proximity.

Segmentation as Clustering - Improvement

- Clustering based on (r,g,b,x,y) values enforces more spatial coherence



Segmentation as Clustering - Challenges



- Color, brightness, position alone are not enough to distinguish all regions...



Segmentation as Clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **texture** similarity




Feature space: filter bank responses (e.g., 24-d)
Kristen Grauman

Cleaning up after segmentation

Problem:

- Histogram-based segmentation can produce messy regions
 - segments do not have to be connected
 - may contain holes



How can these be fixed?

Use Morphological operators!

Image Dilation

Assume:
binary image

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	0	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$F[x, y]$$

1	1	1
1	1	1
1	1	1

$$H[u, v]$$

Dilation: Move mask H over image F , turning F 's pixels to 1 if F and H both have 1s *anywhere* in the region of overlap

- $G[x, y] = 1$ if $H[u, v]$ and $F[x+u-1, y+v-1]$ are both 1 **somewhere**
0 otherwise
- Written as $G = H \oplus F$

Image Dilation

Assume:
binary image

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F[x, y]$


1	1	1
1	1	1
1	1	1

$H[u, v]$

Dilation: Move mask H over image F , turning F 's pixels to 1 if F and H have 1s *anywhere* in the region of overlap

- $G[x, y] = 1$ if $H[u, v]$ and $F[x+u-1, y+v-1]$ are both 1 **somewhere**
0 otherwise
- Written as $G = H \oplus F$

Image Dilation - Example



Demo: <http://www.cs.bris.ac.uk/~majid/mengine/morph.html>

Image Erosion

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	1	0	1	1	1	0	0
0	0	0	1	1	1	1	1	0	0
0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F[x, y]$


1	1	1
1	1	1
1	1	1

$H[u, v]$

Erosion: Move mask H over image F , turning F 's pixels to 1 if F and H both have 1s *everywhere* in the region of overlap

- $G[x, y] = 1$ if $F[x+u-1, y+v-1]$ is 1 **everywhere** that $H[u, v]$ is 1
0 otherwise
- Written $G = H \ominus F$

Image Erosion




Demo: <http://www.cs.bris.ac.uk/~majid/mengine/morph.html>

Nested Dilations and Erosions

What does this operation do?

$$G = H \ominus (H \oplus F)$$



- This is called a **closing** operation

Is this the same thing as the following?

$$G = H \oplus (H \ominus F)$$

Nested Dilations and Erosions

What does this operation do?

$$G = H \oplus (H \ominus F)$$

- This is called an **opening** operation

<http://www.dai.ed.ac.uk/HIPR2/open.htm>

You can clean up binary images (e.g., results of segmentation) by applying combinations of dilations and erosions. Dilations, erosions, opening, and closing operations are known as morphological operations.

see <http://www.dai.ed.ac.uk/HIPR2/morops.htm>

Mean-shift Segmentation

- The mean shift algorithm seeks *modes* or local maxima of density in the feature space

image




Mean-shift Segmentation

- The mean shift algorithm seeks *modes* or local maxima of density in the feature space


image




Feature space
($L^*u^*v^*$ color values)




Mean-Shift



Mean-Shift




Mean-Shift




Mean-Shift Clustering

- Cluster: all data points in the attraction basin of a mode
- Attraction basin: the region for which all trajectories lead to the same mode




Mean-Shift Segmentation


- Find features (color, gradients, texture, etc)
- Initialize windows at individual feature points (pixels)
- Perform mean shift for each window (pixel) until convergence
- Merge windows (pixels) that end up near the same “peak” or mode



(a)



(b)



Mean-Shift Segmentation Results



<http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html>



Mean-Shift Segmentation Results



Mean-Shift Segmentation

- Pros:
 - Does not assume shape on clusters
 - One parameter choice (window size)
 - Generic technique
 - Find multiple modes
- Cons:
 - Selection of window size
 - Does not scale well with dimension of feature space

Images as Graphs




- *Fully-connected* graph
 - node (vertex) for every pixel
 - link between every pair of pixels, p, q
 - affinity weight w_{pq} for each link (edge)
 - w_{pq} measures *similarity*
 - similarity is *inversely proportional* to difference (in color and position...)

Measuring Affinity

- One possibility:



$$\text{aff}(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{1}{2\sigma^2} \text{dist}(\mathbf{x}_i, \mathbf{x}_j)^2\right)$$

Small sigma:
group only
nearby points




Large sigma:
group distant
points

Segmentation by Graph Partitioning



- Break Graph into Segments
 - Delete links that cross between segments
 - Easiest to break links that have low affinity
 - similar pixels should be in the same segments
 - dissimilar pixels should be in different segments

Segmentation by Graph Partitioning




- Set of edges whose removal makes a graph disconnected
- Cost of a cut: sum of weights of cut edges

$$cut(A, B) = \sum_{p \in A, q \in B} w_{p,q}$$

- A graph cut gives us a segmentation
 - What is a “good” graph cut and how do we find one?

Segmentation by Graph Partitioning




- Set of edges whose removal makes a graph disconnected
- Cost of a cut: sum of weights of cut edges

$$cut(A, B) = \sum_{p \in A, q \in B} w_{p,q}$$

- A graph cut gives us a segmentation
 - What is a “good” graph cut and how do we find one?

Min Cut




$$cut(A, B) = \sum_{p \in A, q \in B} w_{p,q}$$

Find minimum cut

- gives you a segmentation

Normalized Cut



Normalized Cut

- fix bias of Min Cut by **normalizing** for size of segments:


$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$$

$assoc(A, V)$ = sum of weights of all edges that touch A

- Ncut value small when we get two clusters with many edges with high weights, and few edges of low weight between them
- Approximate solution for minimizing the Ncut value : generalized eigenvalue problem.

Min Cut - Example

- Minimum Cut




$$Cut(BCDE, A) = 0.17$$

Normalized Cut - Example


$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$$

$assoc(A, V)$ = sum of weights of all edges that touch A




$$NormCut(BCDE, A) = :$$

Normalized Cut - Example



$$NormCut(BCDE, A) = 1.06$$

Normalized Cut – Generalizes better



$$NormCut(BCDE, A) = 1.06$$

$$NormCut(ABC, DE) = 1.03$$

Graph Cut - Results



<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Normalized Cuts – Pros and Cons

Pros:

- Generic framework, flexible to choice of function that computes weights (“affinities”) between nodes
- Does not require model of the data distribution

Cons:

- Time complexity can be high
 - Dense, highly connected graphs → many affinity computations
 - Solving eigenvalue problem
- Preference for balanced partitions

Normalized Cuts – Pros and Cons

Pros:

- Generic framework, flexible to choice of function that computes weights (“affinities”) between nodes
- Does not require model of the data distribution

Cons:

- Time complexity can be high
 - Dense, highly connected graphs → many affinity computations
 - Solving eigenvalue problem
- Preference for balanced partitions