



Image Features and Matching

CSE473/573 - Computer Vision and Image Processing

Junsong Yuan

jsyuan@buffalo.edu

2018 Fall

Image Features

1 Image Features

2 Feature Matching

Image Features - Motivation



Find this book in the image below

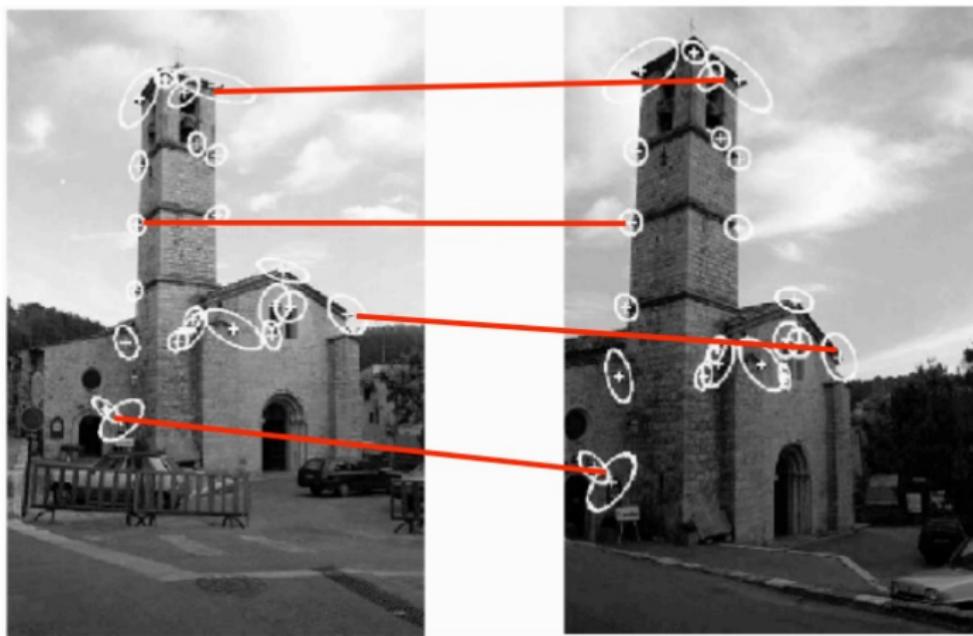


Transformations:

- Orientation
- Scale
- Occlusion

Image Features - Motivation

Vision tasks such as stereo and motion estimation require finding corresponding features across two or more views.



Panorama Creation

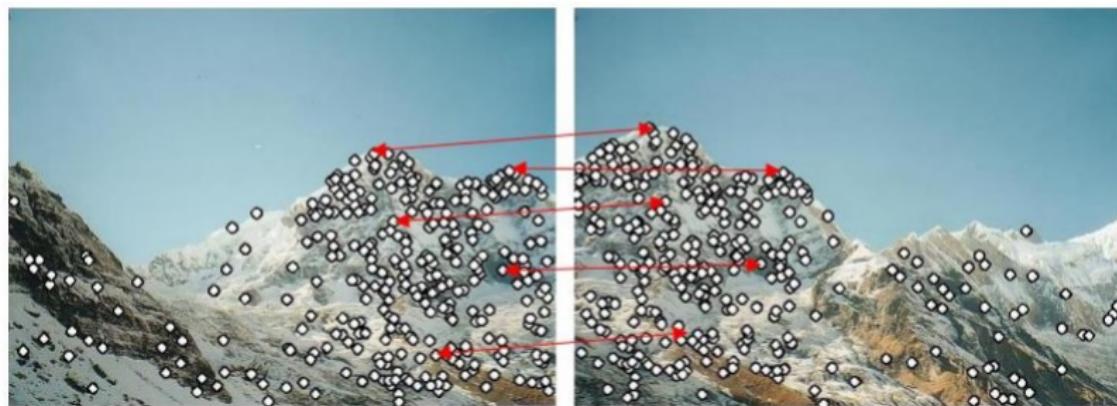
First step is to match and align the images

Find features to do it.



Panorama Creation

Model fitting and RANSAC (Random Sample Consensus)



Panorama Creation

Using geometric transformation



Advantages of Local Features

Locality

- features are local, so robust to occlusion and clutter

Distinctiveness:

- can differentiate a large database of objects

Quantity

- hundreds or thousands in a single image

Efficiency

- real-time performance achievable

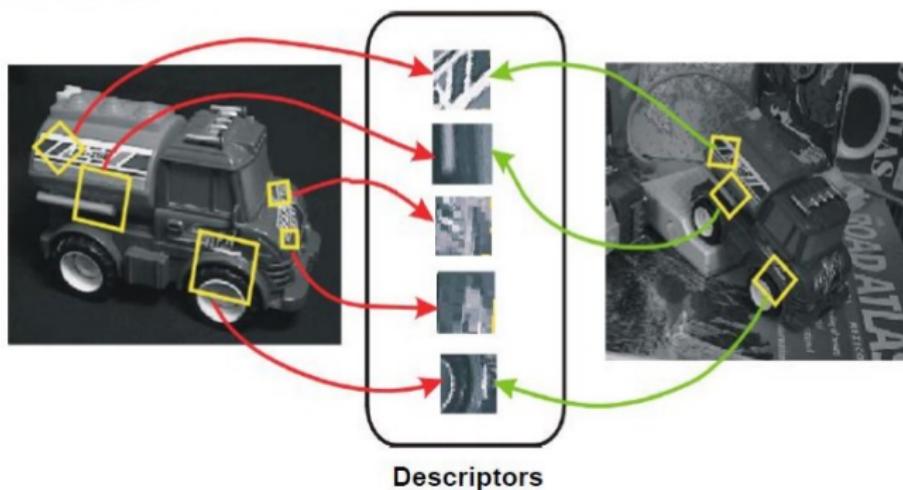
Generality

- exploit different types of features in different situations

Source: S. Seitz slides.

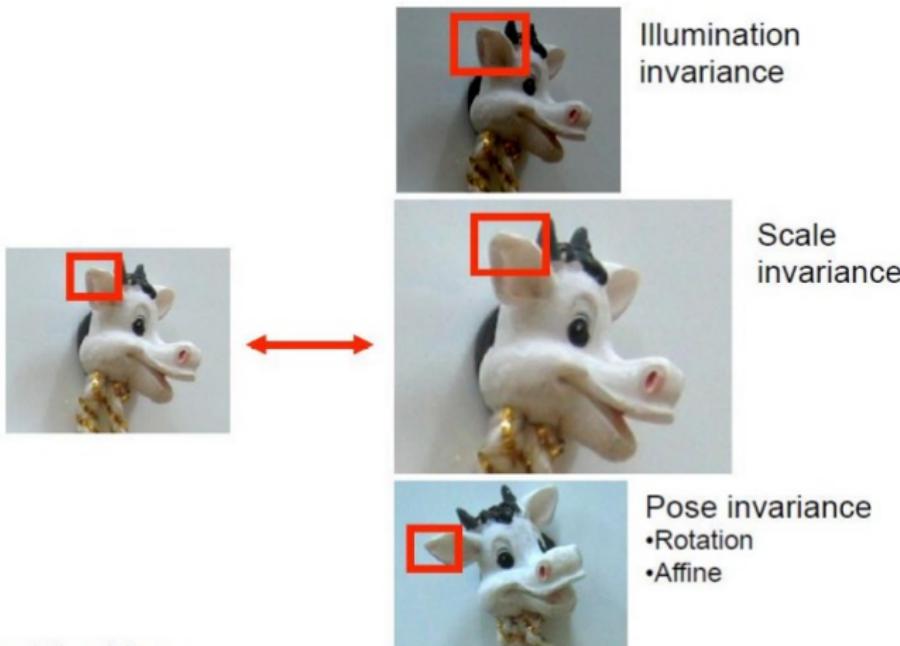
Challenges

- Repeatability
- Uniqueness
- Invariance



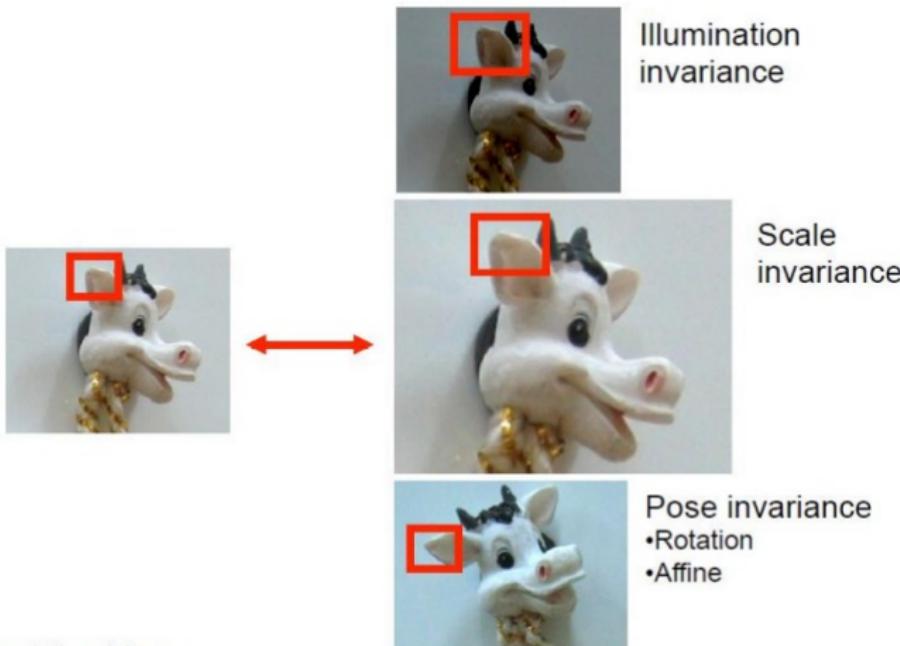
Sources for this example: I. Kokkinos, D. Frolova, D. Simakov, S. Seitz.

Challenges



Source for this example: Savarese.

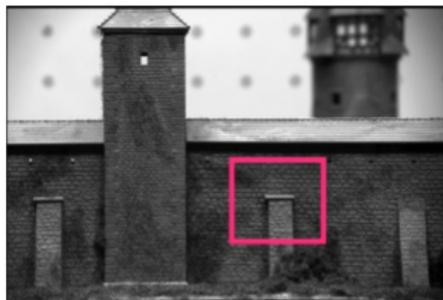
Challenges



Source for this example: Savarese.

What features would you pick?

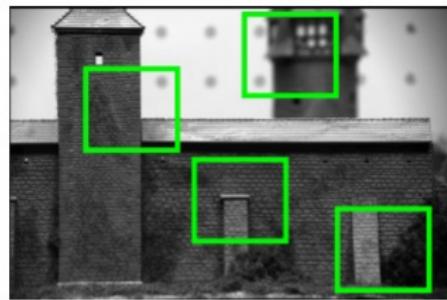
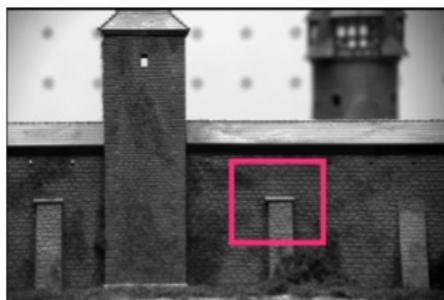
Elements to be matched are image patches of fixed size



Task: find the best (most similar) patch in a second image



What features would you pick?



Intuition: this would be a good patch for matching, since it is very distinctive (there is only one patch in the second frame that looks similar).



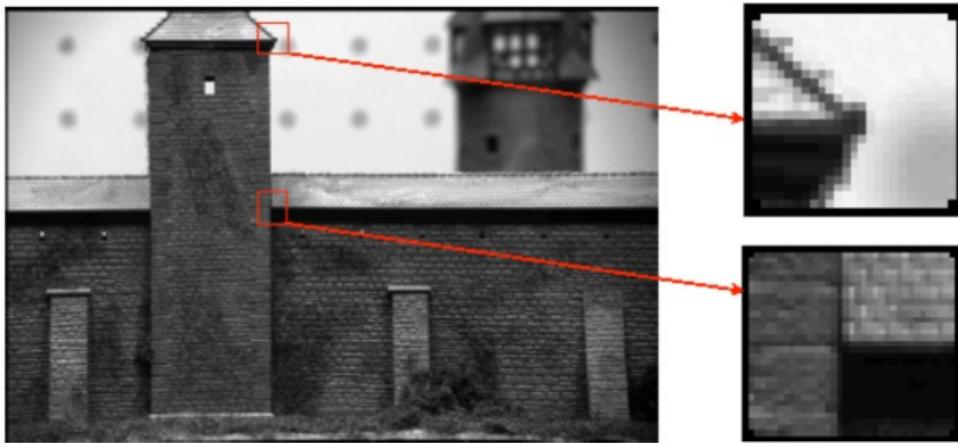
What features would you pick?



Intuition: this would be a BAD patch for matching, since it is not very distinctive (there are many similar patches in the second frame)



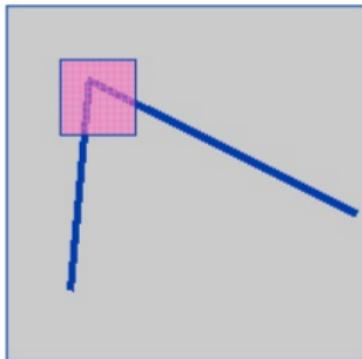
Corners as Features



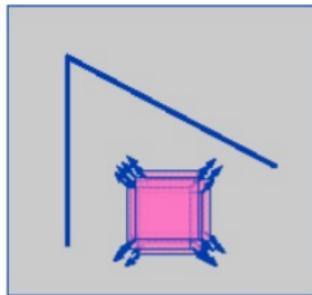
- Intuitively, junctions of contours.
- Generally more stable features over changes of viewpoint
- Intuitively, large variations in the neighborhood of the point in all directions
- They are good features to match!

Corners as Features

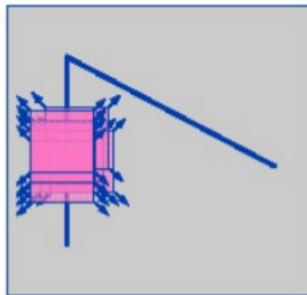
- We should easily recognize the point by looking at intensity values within a small window
- Shifting the window in *any direction* should yield a *large change* in appearance.



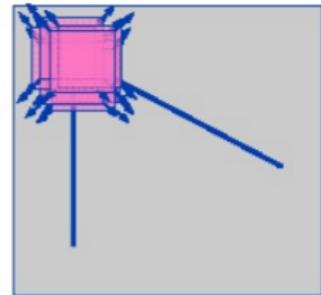
Harris Corner Detector



“flat” region:
no change in
all directions



“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions

Harris corner detector gives a mathematical approach for determining which case holds.

Harris Detector - Math

Change of intensity for the shift $[u, v]$:

$$E(u, v) = \sum_{x, y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

The diagram illustrates the Harris detector equation. It shows the equation $E(u, v) = \sum_{x, y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$ enclosed in a green box. Below the box, three green ovals are connected by arrows to the terms in the equation: 'Window function' points to $w(x, y)$, 'Shifted intensity' points to $I(x+u, y+v)$, and 'Intensity' points to $I(x, y)$.

For nearly constant patches, this will be near 0.
For very distinctive patches, this will be larger.
Hence... we want patches where $E(u, v)$ is LARGE.

Taylor Series

$$f(x+u, y+v) = f(x, y) + u f_x(x, y) + v f_y(x, y) +$$

First partial derivatives

$$\frac{1}{2!} [u^2 f_{xx}(x, y) + uv f_{xy}(x, y) + v^2 f_{yy}(x, y)] +$$

Second partial derivatives

$$\frac{1}{3!} [u^3 f_{xxx}(x, y) + u^2 v f_{xxy}(x, y) + u v^2 f_{xyy}(x, y) + v^3 f_{yyy}(x, y)]$$

Third partial derivatives

+ ... (Higher order terms)

First order approx

$$f(x+u, y+v) \approx f(x, y) + u f_x(x, y) + v f_y(x, y)$$

Harris Corner Derivation

$$\begin{aligned} & \sum [I(x+u, y+v) - I(x, y)]^2 \\ & \approx \sum [I(x, y) + uI_x + vI_y - I(x, y)]^2 \quad \text{First order approx} \\ & = \sum u^2 I_x^2 + 2uvI_xI_y + v^2 I_y^2 \\ & = \sum \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} I_x^2 & I_xI_y \\ I_xI_y & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} \quad \text{Rewrite as matrix equation} \\ & = \begin{bmatrix} u & v \end{bmatrix} \left(\sum \begin{bmatrix} I_x^2 & I_xI_y \\ I_xI_y & I_y^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

Harris Corner Derivation

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Widening function - computing a weighted sum (simplest case, $w=1$)

Note: these are just products of components of the gradient, I_x, I_y

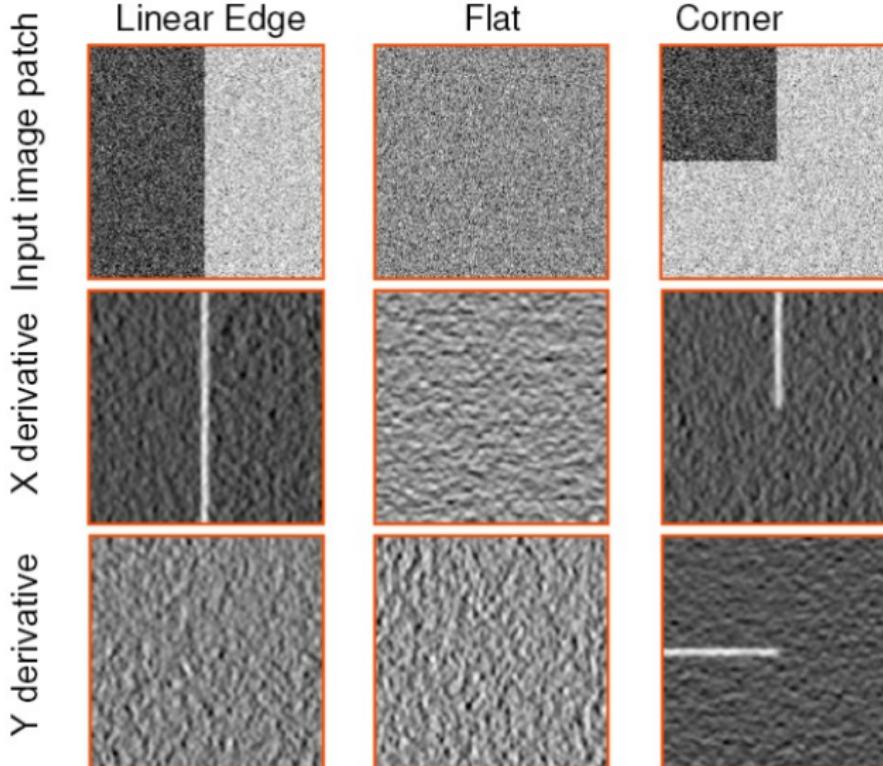
Harris Detector - Intuition

Treat gradient vectors as a set of (dx, dy) points with a center of mass defined as being at $(0,0)$.

Fit an ellipse to that set of points via scatter matrix

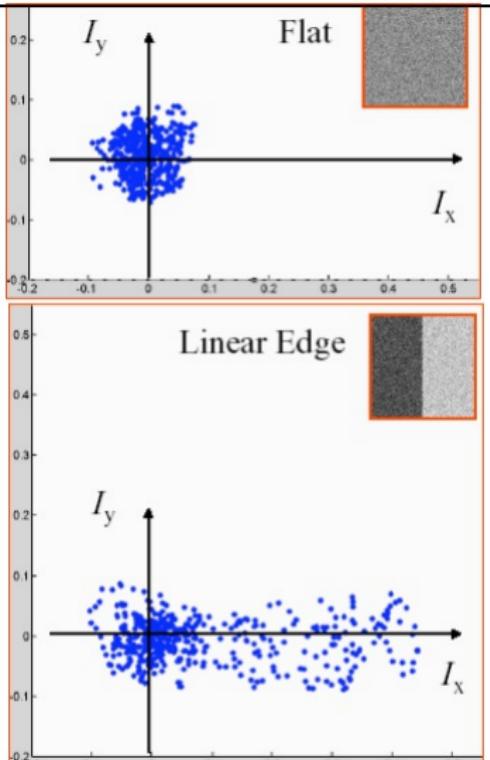
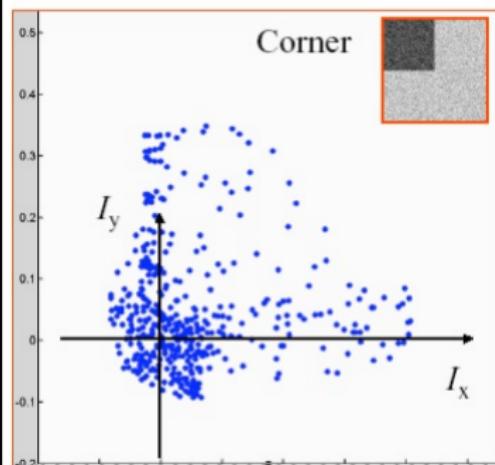
Analyze ellipse parameters for varying cases...

Example Images



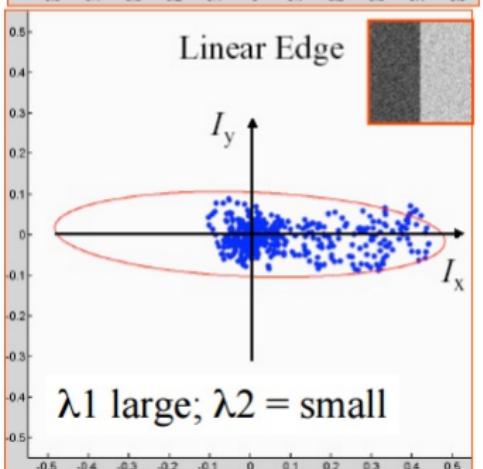
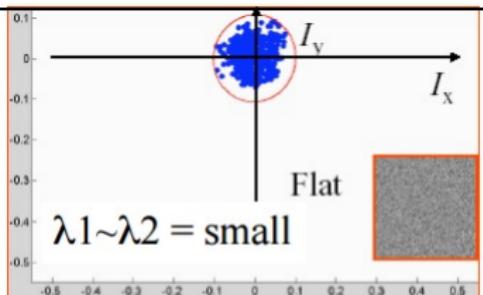
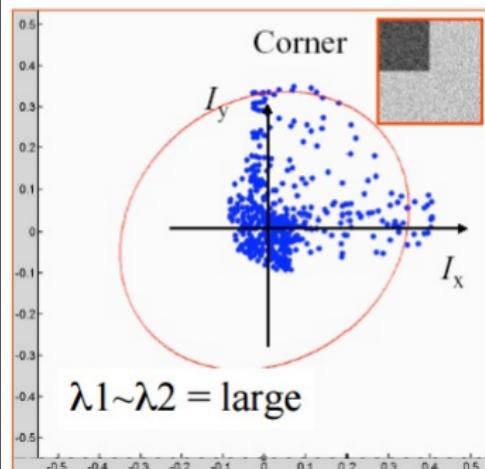
Plotting the gradients

The distribution of the x and y derivatives is very different for all three types of patches



Fitting Ellipses to the set of points

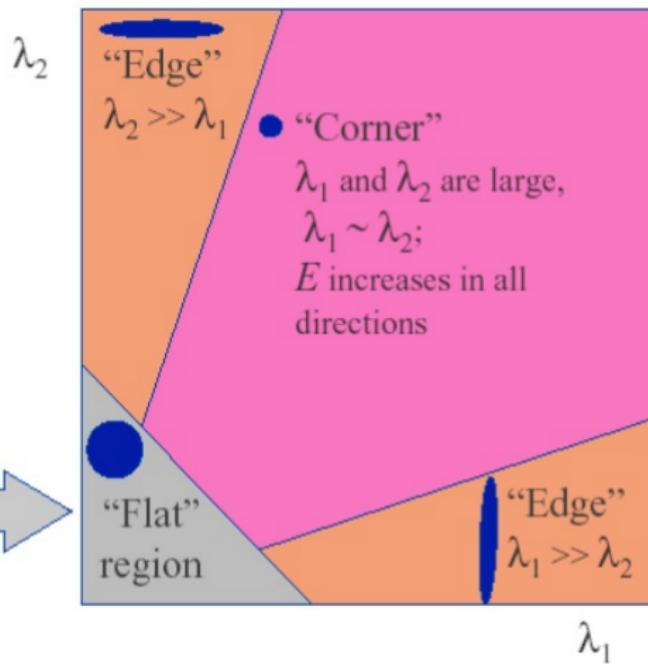
The distribution of x and y derivatives can be characterized by the shape and size of the principal component ellipse



Classification via Eigen Values

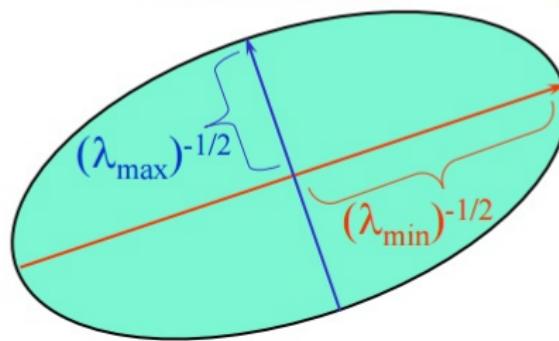
Classification of image points using eigenvalues of M :

λ_1 and λ_2 are small;
 E is almost constant in all directions



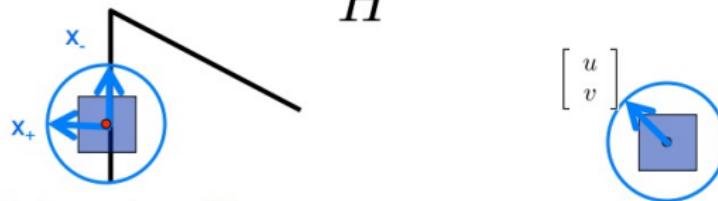
Classification via Eigen Values

$$E(u, v) \equiv [u, v] M \begin{bmatrix} u \\ v \end{bmatrix} \quad \lambda_1, \lambda_2 - \text{eigenvalues of } M$$

Ellipse $E(u, v) = \text{const}$ direction of the
fastest changedirection of the
slowest change

Classification via Eigen Values

$$E(u, v) = \sum_{(x,y) \in W} [u \ v] \underbrace{\begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

Eigenvalues and eigenvectors of H

- Define shifts with the smallest and largest change (E value)
- x_+ = direction of largest increase in E.
- λ_+ = amount of increase in direction x_+
- x_- = direction of smallest increase in E.
- λ_- = amount of increase in direction x_-

$$\begin{aligned} Hx_+ &= \lambda_+ x_+ \\ Hx_- &= \lambda_- x_- \end{aligned}$$

Source: S. Seitz.

Eigen Values Computation

The **eigenvectors** of a matrix \mathbf{A} are the vectors \mathbf{x} that satisfy:

$$\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$$

The scalar λ is the **eigenvalue** corresponding to \mathbf{x}

- The eigenvalues are found by solving:

$$\det(\mathbf{A} - \lambda\mathbf{I}) = 0$$

- In our case, $\mathbf{A} = \mathbf{H}$ is a 2×2 matrix, so we have

$$\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$$

- The solution:

$$\lambda_{\pm} = \frac{1}{2} \left[(h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$$

Once you know λ , you find \mathbf{x} by solving

$$\begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$

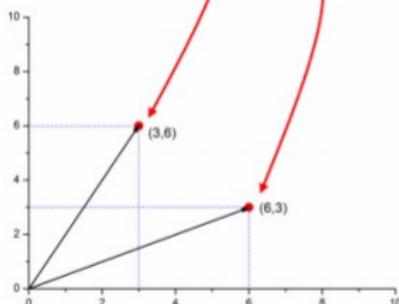
Source: S. Seitz.

Eigen Values Numerical Example

Numerical Example 1

$$[A] = \begin{bmatrix} 3 & 6 \\ 6 & 3 \end{bmatrix}$$

Let's get the eigenvalues of matrix [A]



$$\begin{vmatrix} 3-\lambda & 6 \\ 6 & 3-\lambda \end{vmatrix} = 0$$

$$(3-\lambda)^2 - 6 \times 6 = 0$$

$$9 - 6\lambda + \lambda^2 - 36 = 0$$

$$\lambda^2 - 6\lambda - 27 = 0$$

$$(\lambda - 9)(\lambda + 3) = 0$$

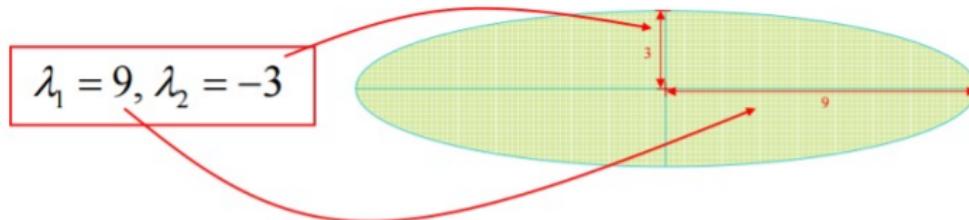
$$\lambda_1 = 9, \lambda_2 = -3$$

Two eigenvalues from
2 variables

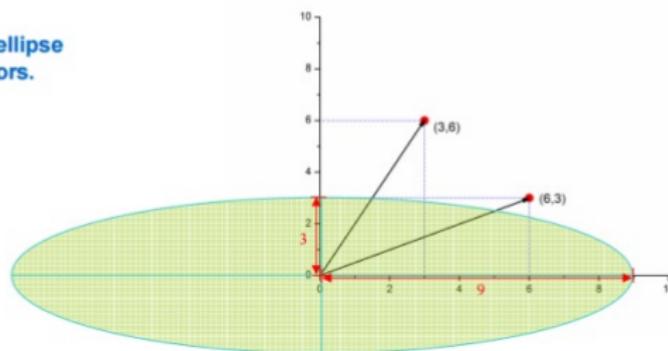
The eigenvalues represent the magnitudes, or lengths, of the major and minor axes of an ellipse.
See Next Page!

Eigen Values Numerical Example

The eigenvalues represent the magnitudes, or lengths, of the major and minor axes of an ellipse.

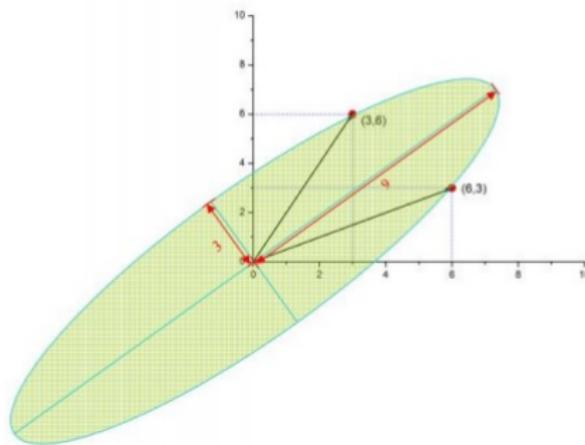


Let's overlap the center of ellipse onto center of the two vectors.



Eigen Values Numerical Example

Then rotate the ellipse to its envelope can be on the two points.



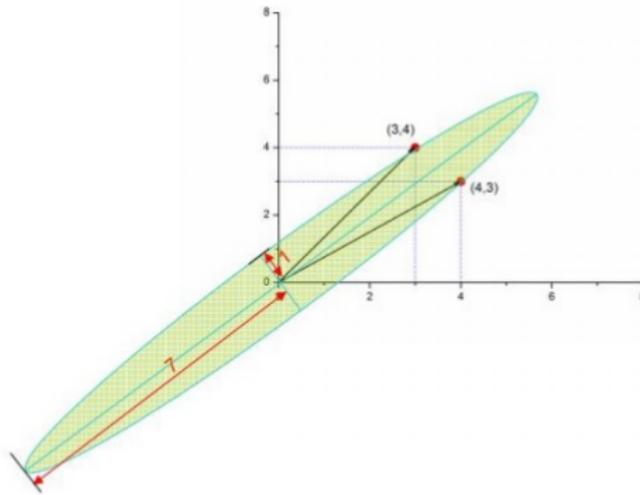
Eigen Values Numerical Example

If the two points are closer than the case of numerical example 1,

$$[A] = \begin{bmatrix} 3 & 4 \\ 4 & 3 \end{bmatrix}$$

Then the eigenvalues are

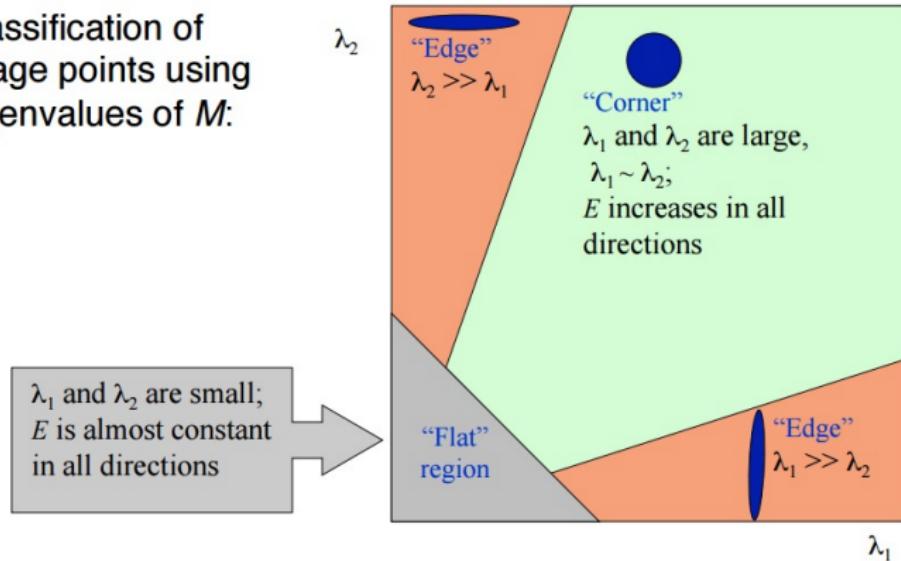
$$\lambda_1 = 7, \lambda_2 = -1$$



The sum of the eigenvalues ($=\lambda_1 + \lambda_2 = 7 - 1 = 6$) is always equal to the sum of the diagonal elements (3+3), i.e. trace of the original matrix.

Classification via Eigen Values

Classification of image points using eigenvalues of M :



Source: Kokkinos, Savarese.

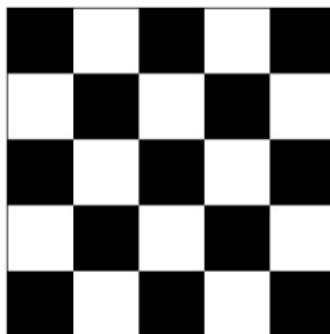
Classification via Eigen Values

How are λ_+ , x_+ , λ_- , and x_- relevant for feature detection?

- What's our feature scoring function?

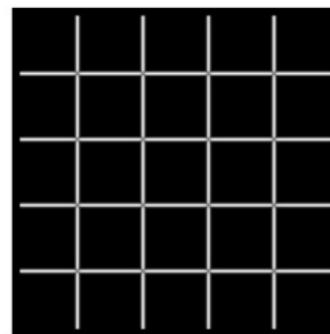
Want $E(u,v)$ to be *large* for small shifts in *all* directions

- the *minimum* of $E(u,v)$ should be large, over all unit vectors $[u \ v]$
- this minimum is given by the smaller eigenvalue (λ_-) of H

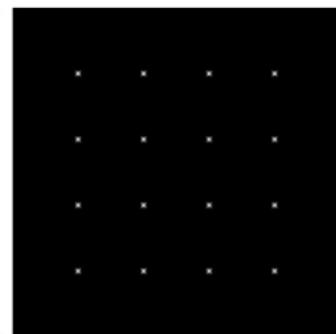


Source: S. Seitz.

I



λ_+

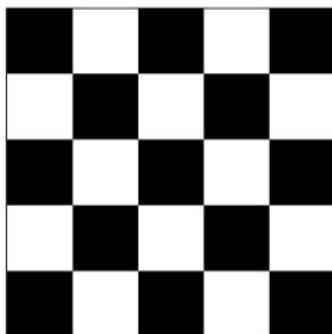


λ_-

Corner Detection Pipeline

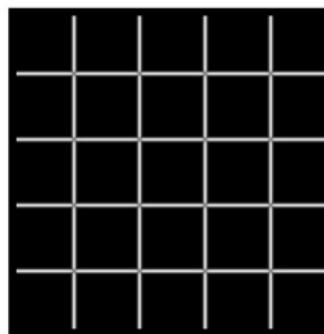
Here's what you do

- Compute the gradient at each point in the image
- Create the H matrix from the entries in the gradient
- Compute the eigenvalues.
- Find points with large response ($\lambda_- > \text{threshold}$)
- Choose those points where λ_- is a local maximum as features

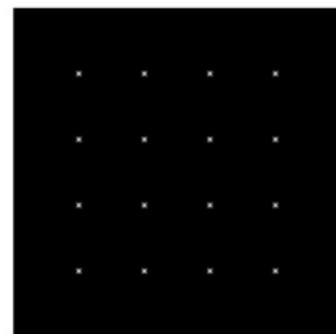


I

Source: S. Seitz.



λ_+



λ_-

Harris Operator

λ_c is a variant of the “Harris operator” for feature detection

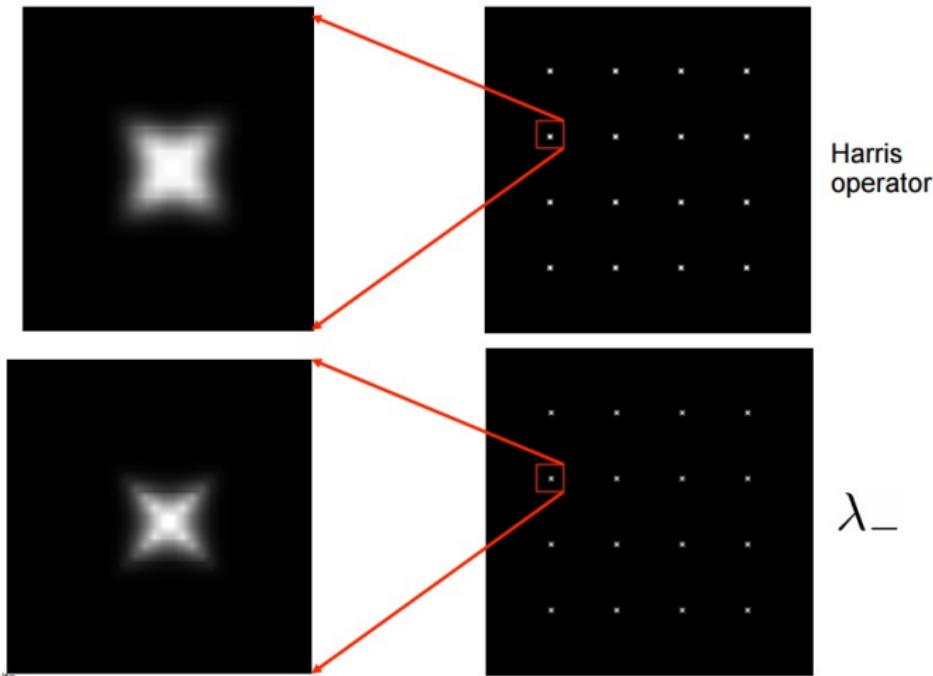
$$\begin{aligned}f &= \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2} \\&= \frac{\text{determinant}(H)}{\text{trace}(H)}\end{aligned}$$

- The *trace* is the sum of the diagonals, i.e., $\text{trace}(H) = h_{11} + h_{22}$
- Very similar to λ_c , but less expensive (no square root)
- Called the “Harris Corner Detector” or “Harris Operator”
- Lots of other detectors, this is one of the most popular

C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#)
Proceedings of the 4th Alvey Vision Conference: pages 147--151. 1988.

Source: S. Seitz.

Harris Operator

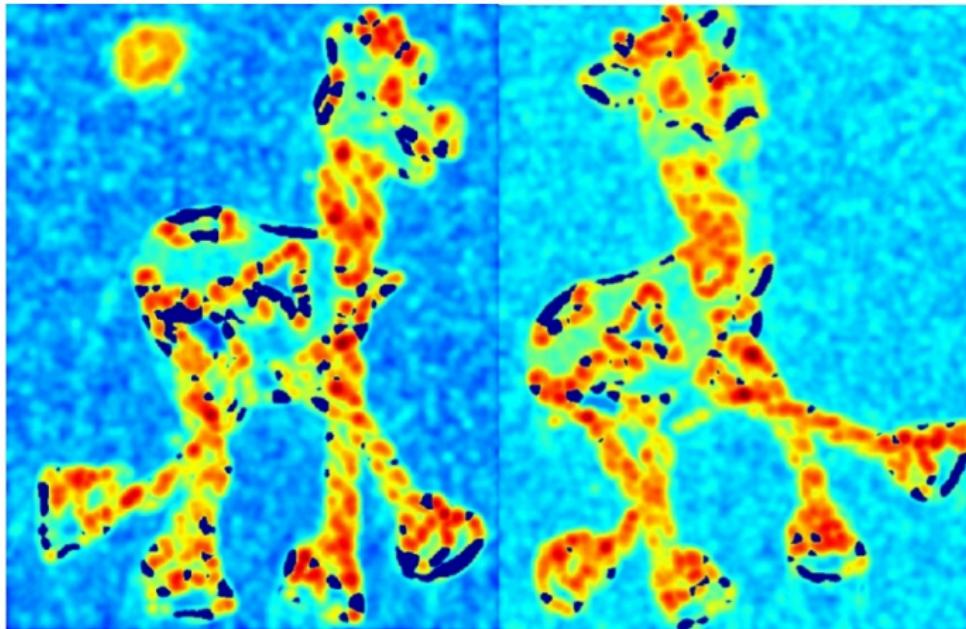


Source: S. Seltz.

Corner Detection Example

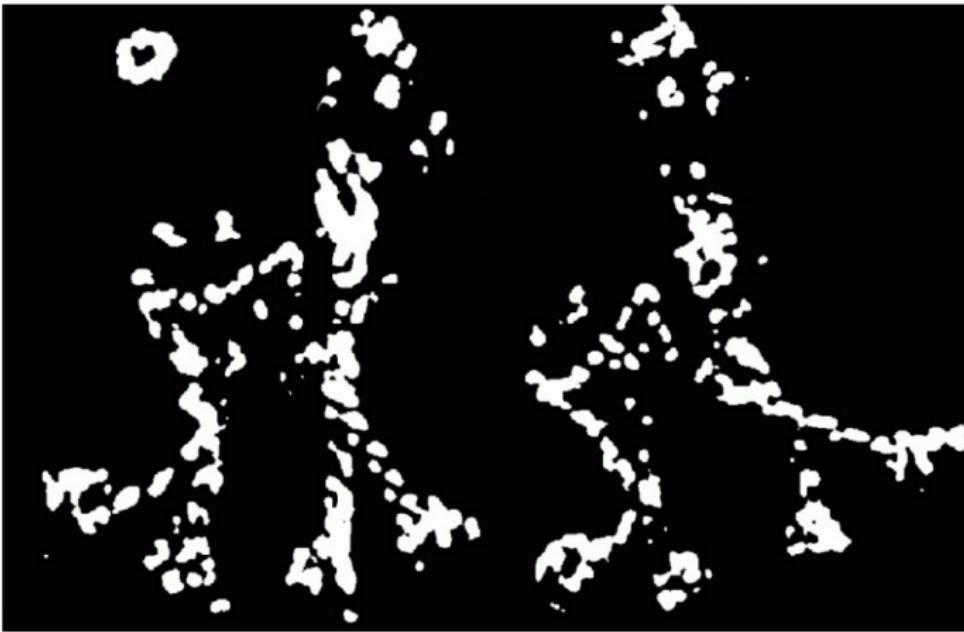


Corner Detection Example



Source: S. Seitz.

Corner Detection Example



Source: S. Seitz.

Corner Detection Example

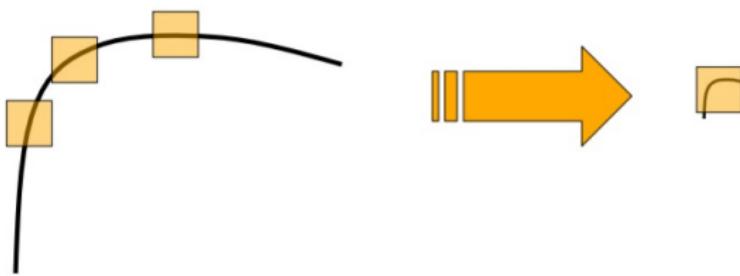


Source: S. Seitz.

Corner Detection Example



Are Corners Scale Invariant?

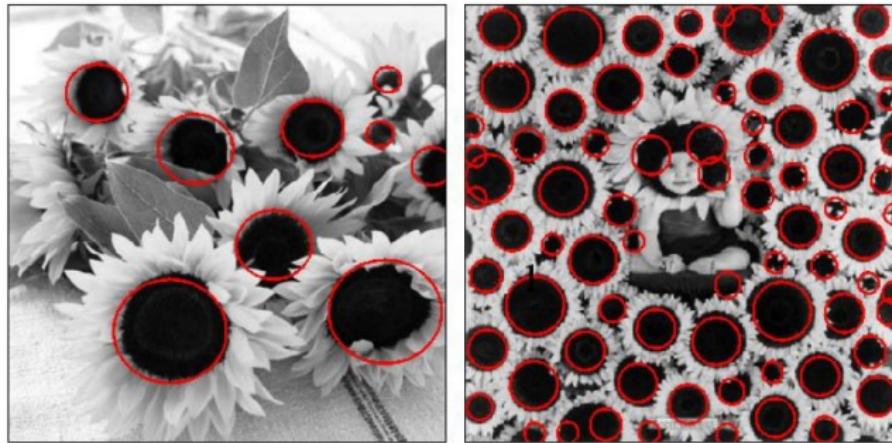


All points will be
classified as edges

Corner !

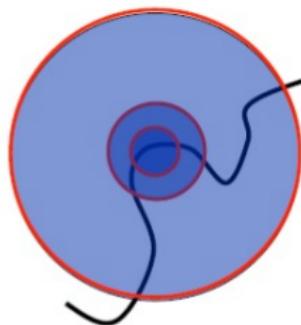
Scale Invariant Features - Blobs

Simple technique to detect features that are scale invariant



Scale Invariant Detection

Suppose you're looking for corners



Key idea: find scale that gives local maximum of f

- f is a local maximum in both position and scale
- Common definition of f : Laplacian
(or difference between two Gaussian filtered images with different sigmas)

Efficiency - Use of DoG vs Laplacian

Computing Difference of Gaussians is simpler!

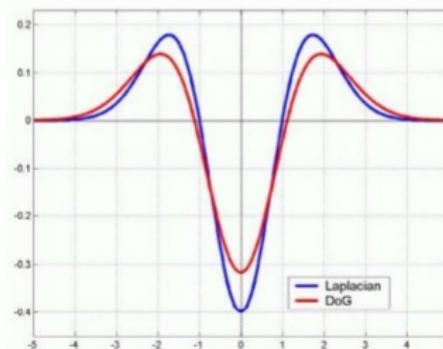
Approximating the Laplacian with a difference of Gaussians:

$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma))$$

(Laplacian)

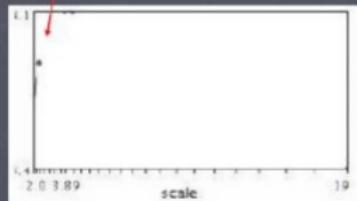
$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



Automatic Scale Selection

Lindeberg et al., 1996



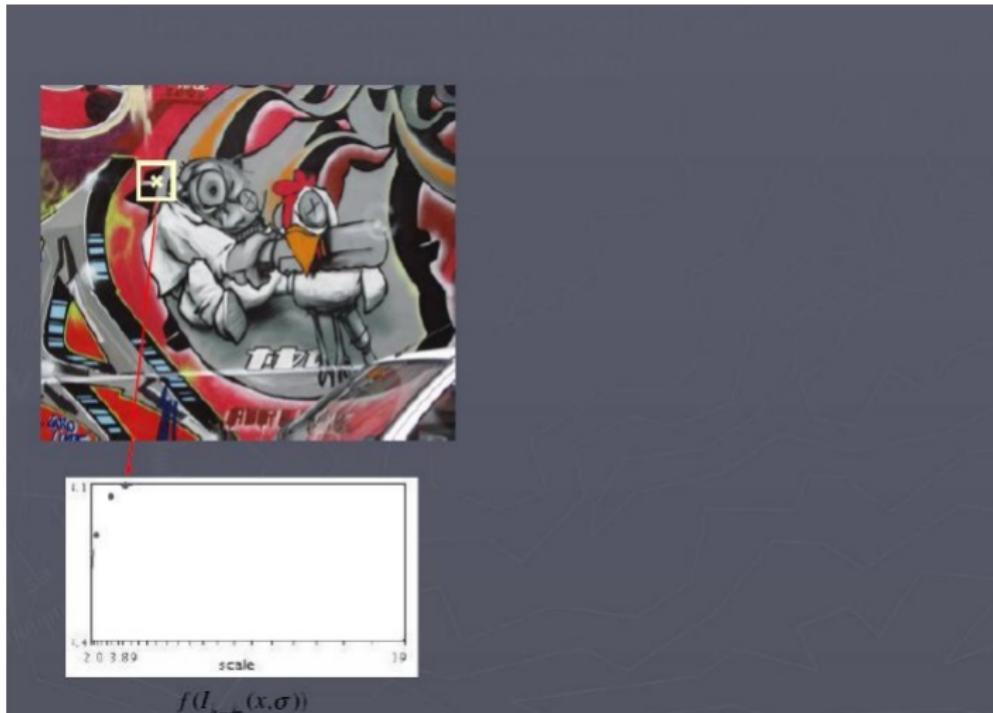
$$f(I_{i-\Delta i}(x, \sigma))$$

Slide from Tinne Tuytelaars

Automatic Scale Selection



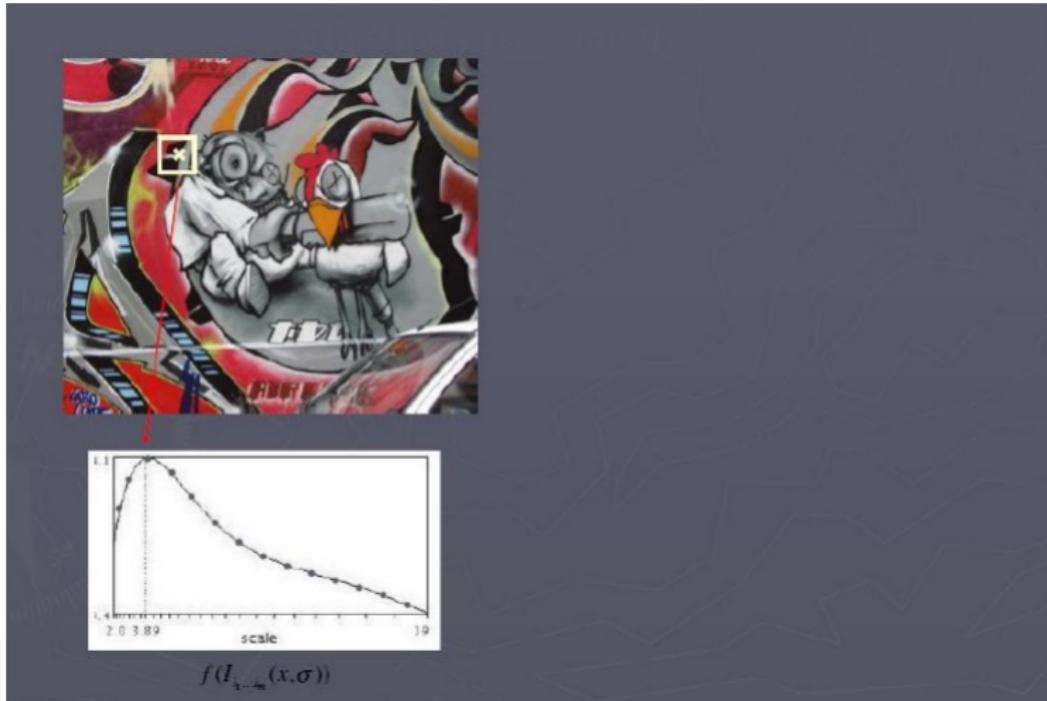
Automatic Scale Selection



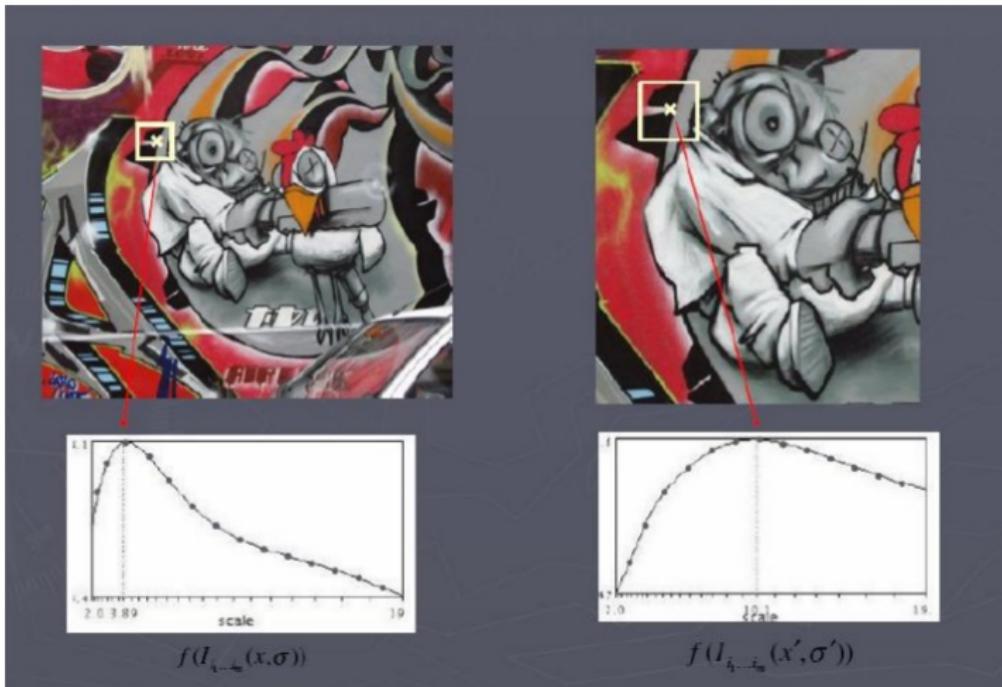
Automatic Scale Selection



Automatic Scale Selection



Automatic Scale Selection



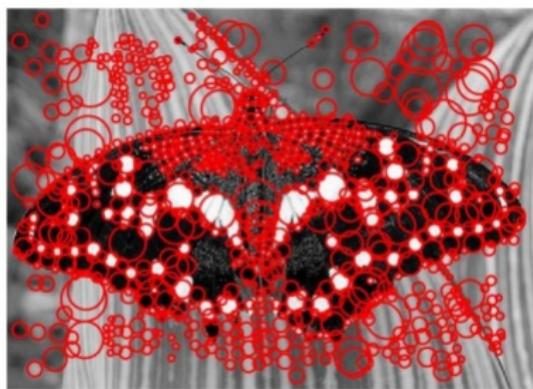
Affine Transformation

Perform Eigen decomposition on the blob window

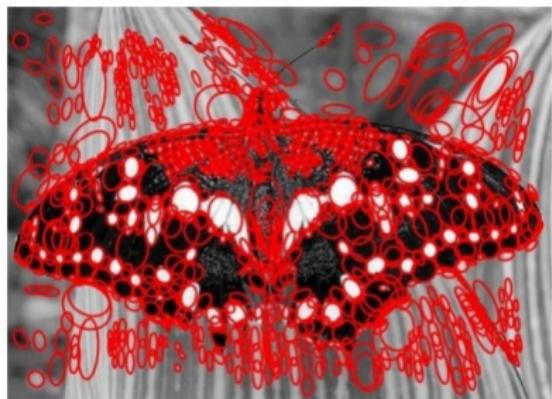
$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$



After Affine Adaptation



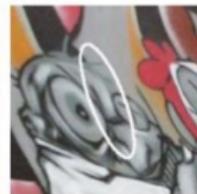
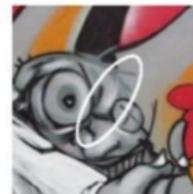
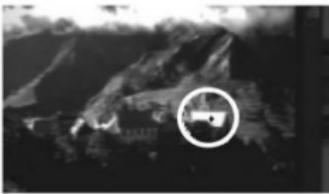
Scale-invariant regions (blobs)



Affine-adapted blobs

Towards Invariant Description

- Geometrically transformed versions of the same neighborhood will give rise to regions that are related by the same transformation
- What to do if we want to compare the appearance of these image regions?
 - *Normalization*: transform these regions into same-size circles



Affine Normalization

Problem: There is no unique transformation from an ellipse to a unit circle

- We can rotate or flip a unit circle, and it still stays a unit circle

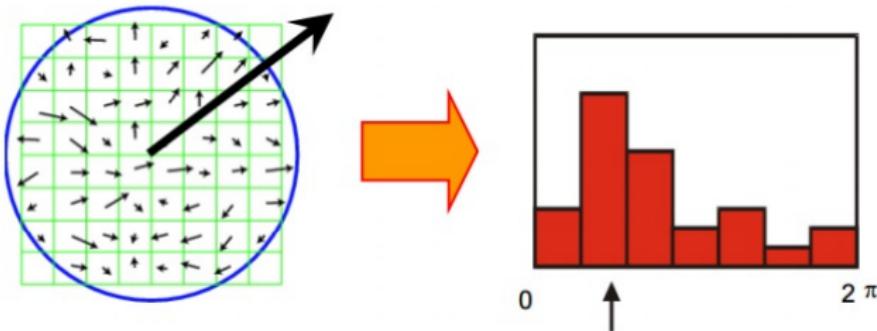


Solution - Gradient Orientations

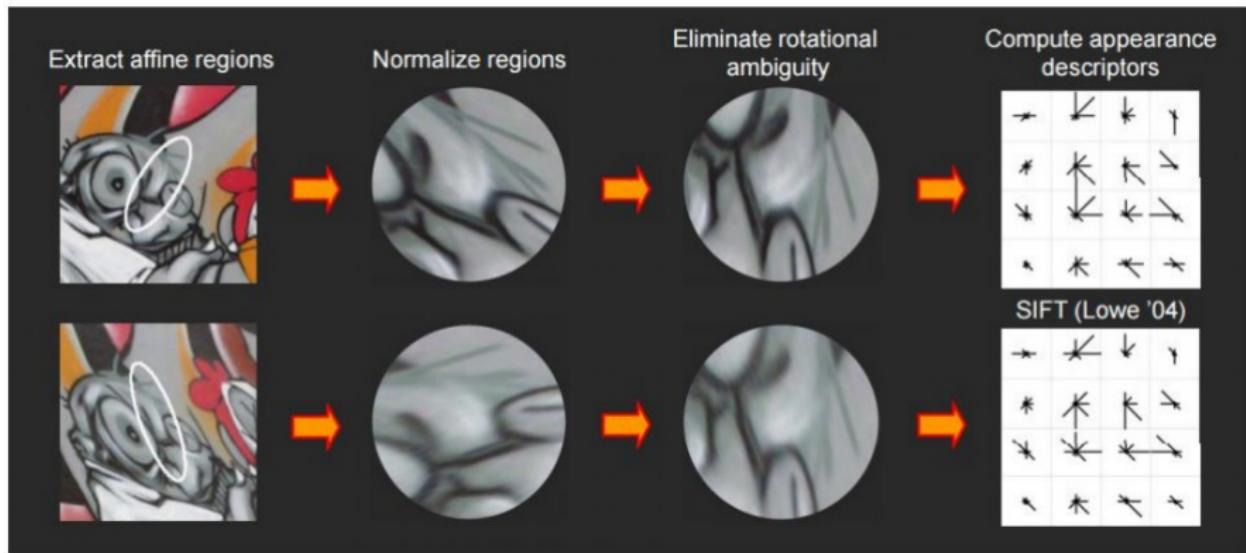


Eliminating Rotational Ambiguity

- To assign a unique orientation to circular image windows:
 - Create histogram of local gradient directions in the patch
 - Assign canonical orientation at peak of smoothed histogram



Scale and Rotational Invariant Features

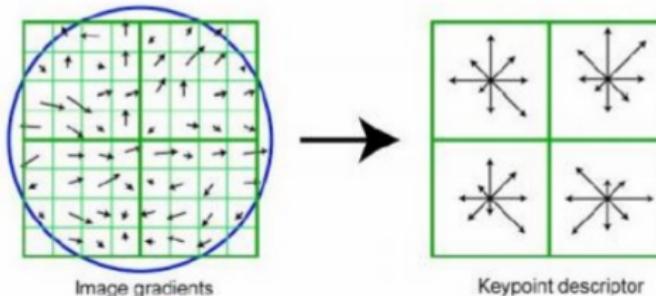


Scale Invariant Feature Transform

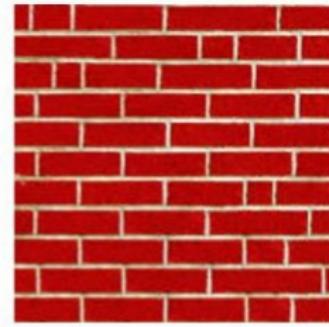
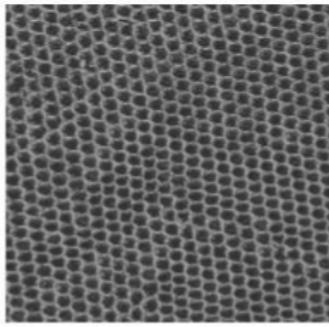
Robert Collins
CSE486, Penn State

SIFT Vector

- Thresholded image gradients are sampled over 16x16 array of locations in scale space
- Create array of orientation histograms
- 8 orientations x 4x4 histogram array = 128 dimensions

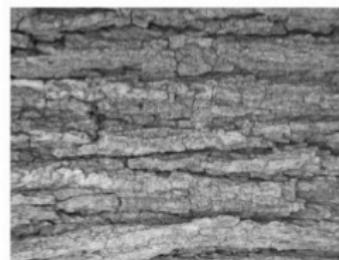
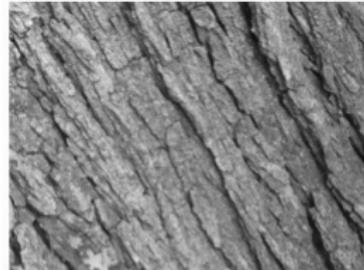


Texture

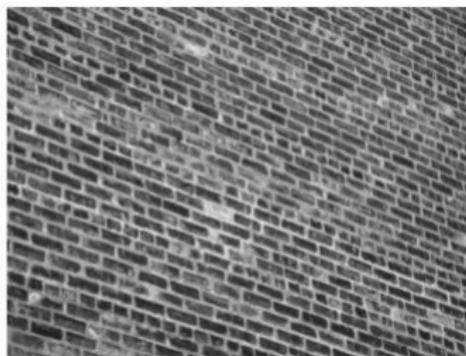
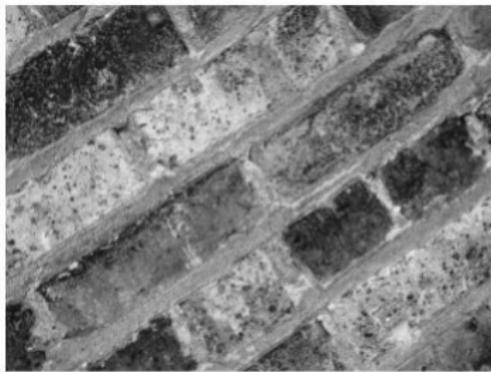


- An image obeying some statistical properties
- Similar structures repeated over and over again
- Often has some degree of randomness

Texture and Orientation



Texture and Scale



Detecting Texture using Spatial Filters

- Look for the subelements
- But what are the subelements, and how do we find them?
- Find subelements by applying filters, looking at the magnitude of the response
- Spots and bars detectors at various scales and orientations.

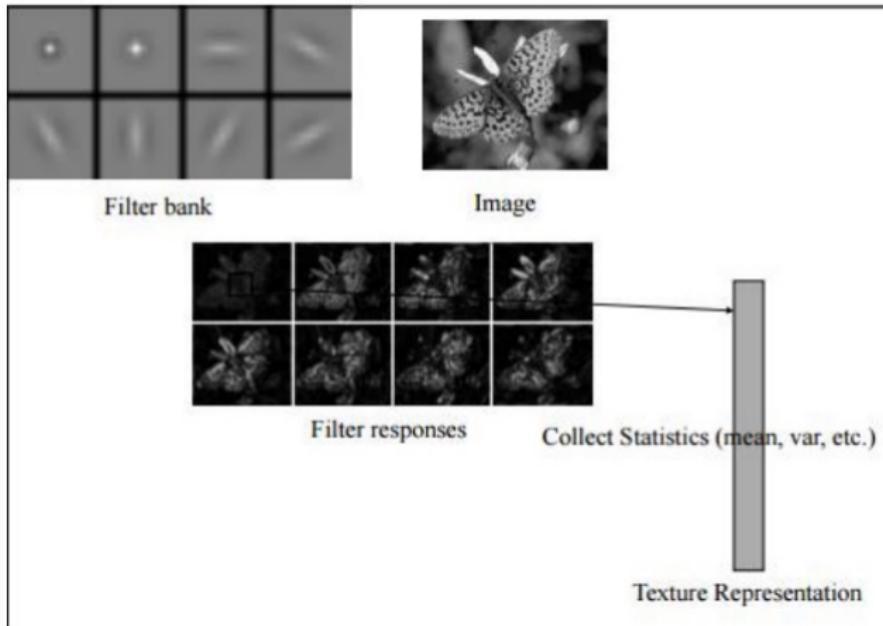
Typically:

- “Spot” filters are Gaussians or weighted sums of concentric Gaussians.
- “Bar” filters are differentiating oriented Gaussians



CS 534 - Texture - 13

Texture - Feature Descriptor



Feature Matching



NASA Mars Rover images
with SIFT feature matches
Figure by Noah Snavely

Feature Matching - Distance Measure

Feature matching

Given a feature in I_1 , how to find the best match in I_2 ?

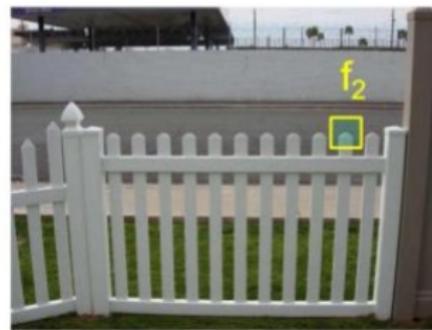
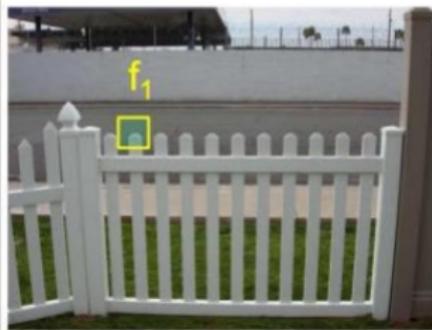
1. Define distance function that compares two descriptors
2. Test all the features in I_2 , find the one with min distance

Feature Matching - L2 Norm

Feature distance

How to define the difference between two features f_1, f_2 ?

- Simple approach: L₂ distance, $\|f_1 - f_2\|$
- can give good scores to ambiguous (incorrect) matches

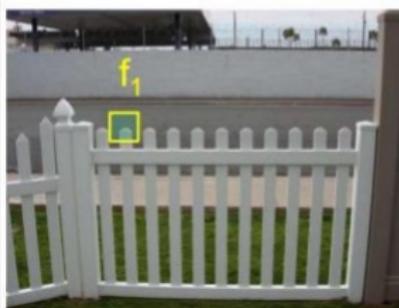


Feature Matching - Ratio of Distance

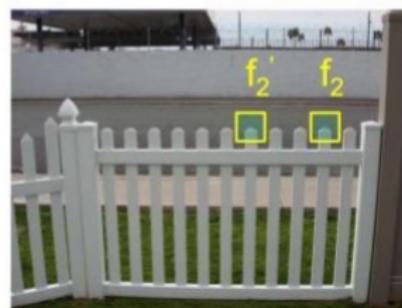
Feature distance

How to define the difference between two features f_1, f_2 ?

- Better approach: ratio distance = $\|f_1 - f_2\| / \|f_1 - f_2'\|$



I_1

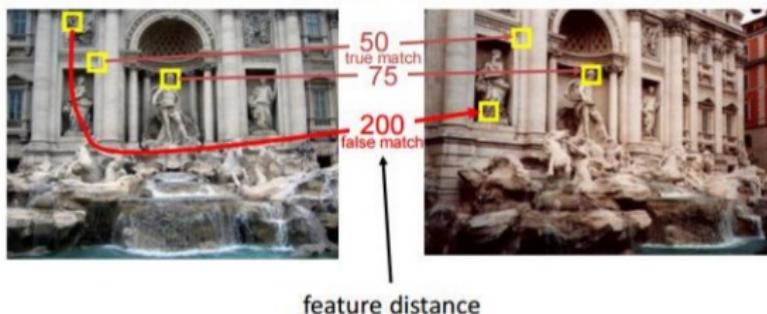


I_2

Feature Matching - Performance

True/false positives

How can we measure the performance of a feature matcher?



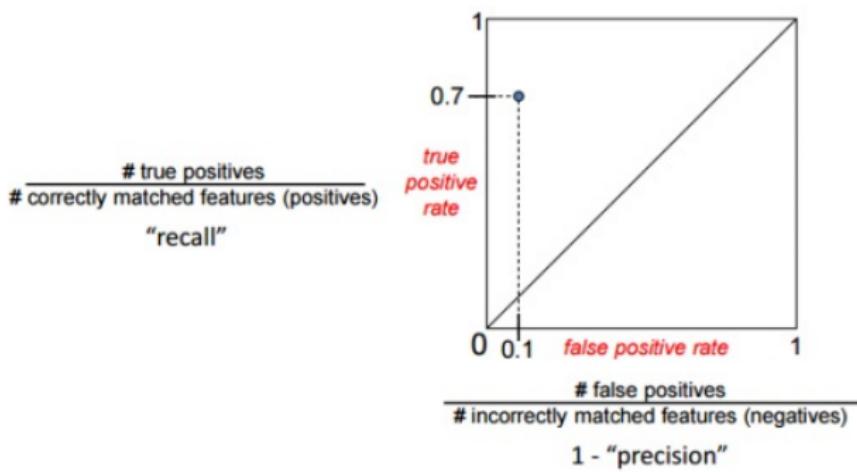
The distance threshold affects performance

- True positives = # of detected matches that are correct
 - Suppose we want to maximize these—how to choose threshold?
- False positives = # of detected matches that are incorrect
 - Suppose we want to minimize these—how to choose threshold?

Feature Matching - Performance

Evaluating the results

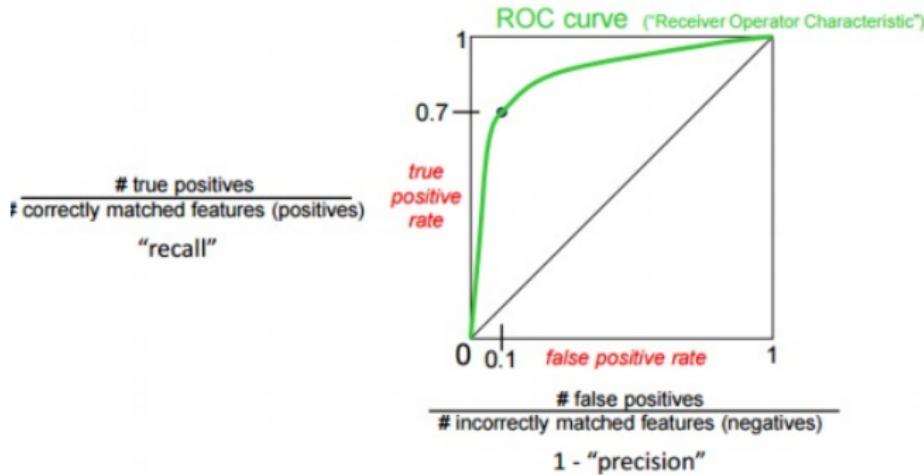
How can we measure the performance of a feature matcher?



Performance Analysis - ROC Curve

Evaluating the results

How can we measure the performance of a feature matcher?

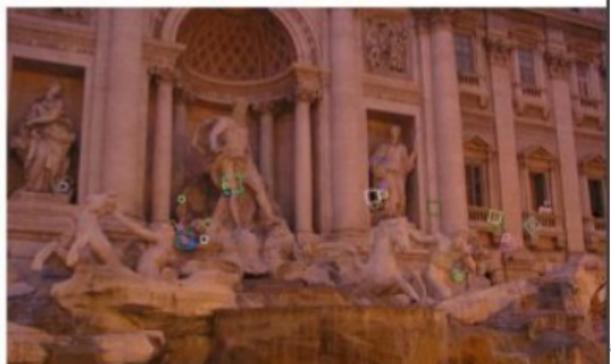


Properties of SIFT based Matching

Extraordinarily robust matching technique

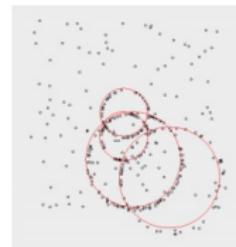
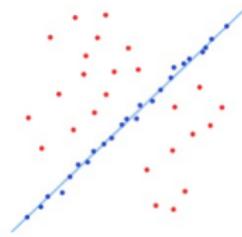
- Can handle **changes in viewpoint**
 - Up to about 60 degree out of plane rotation
- Can handle **significant changes in illumination**: Sometimes even day vs. night (below)
- **Fast and efficient** — can run in real time
- Lots of code available:

http://people.csail.mit.edu/albert/ladypack/wiki/index.php/Known_implementations_of_SIFT



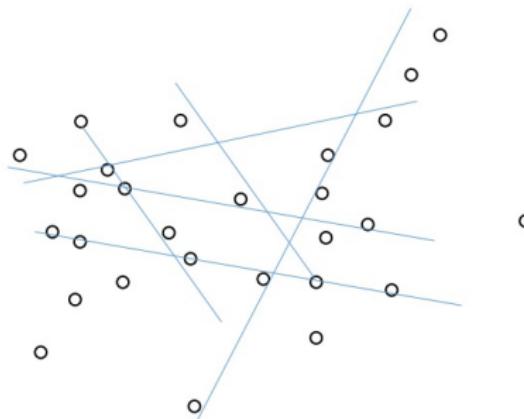
RANSAC

- Random Sample Consensus
- Used for Parametric Matching/Model Fitting
- Applications:



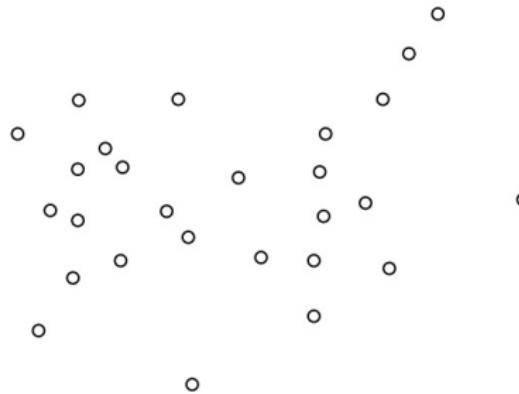
How RANSAC Works

- Fit the best possible Line to these points
- Brute Force Search ?
- Not Feasible
- Better Strategy?



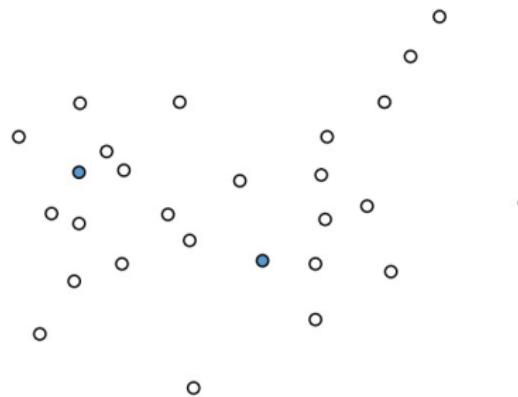
Line Fitting Using RANSAC

- Random Search – Much Faster!!!



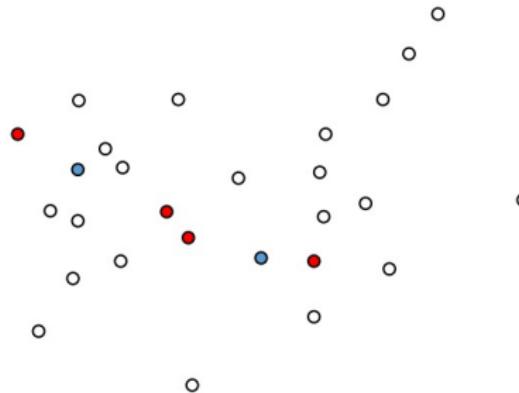
Line Fitting Using RANSAC

- Iteration 1



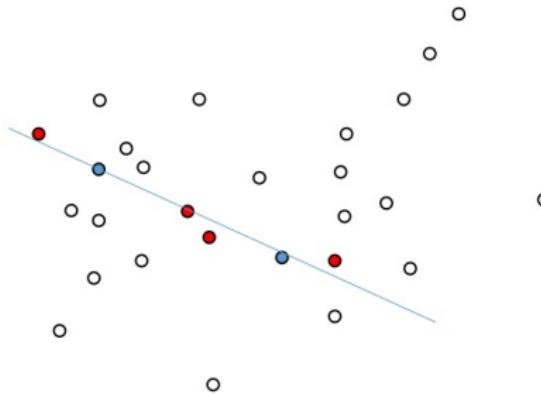
Line Fitting Using RANSAC

- Iteration 1



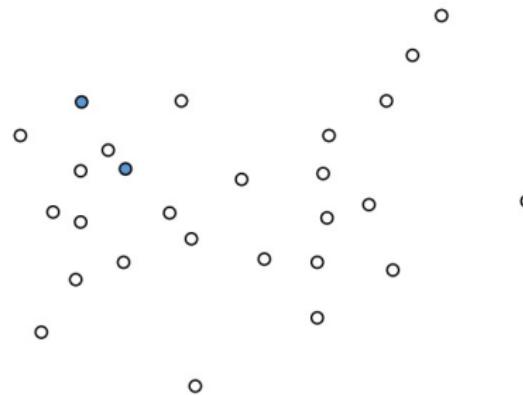
Line Fitting Using RANSAC

- Iteration 1



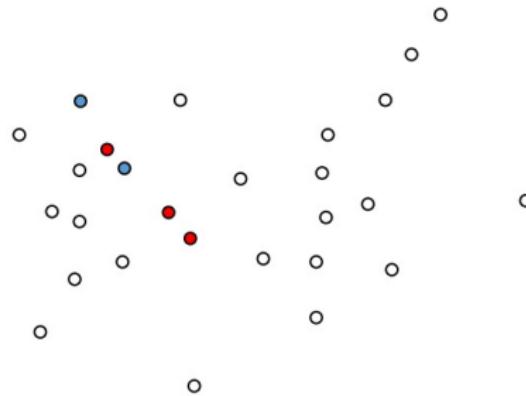
Line Fitting Using RANSAC

- Iteration 2



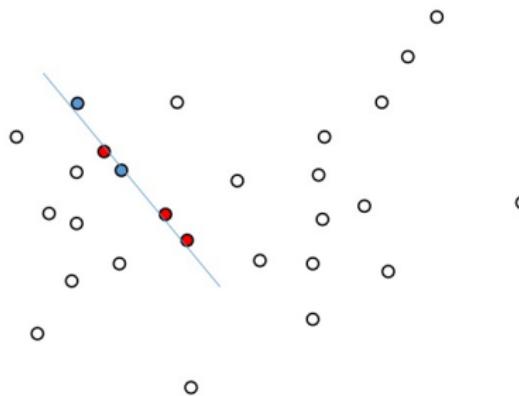
Line Fitting Using RANSAC

- Iteration 2



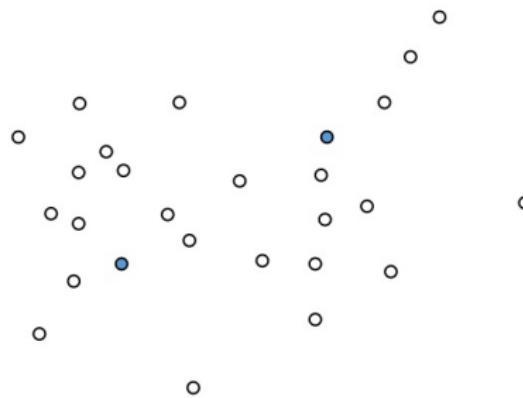
Line Fitting Using RANSAC

- Iteration 2



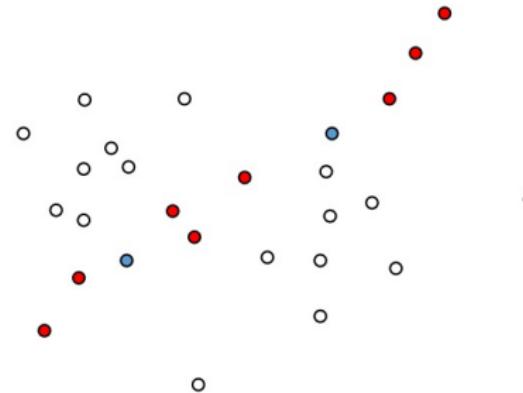
Line Fitting Using RANSAC

- ...
- Iteration 5



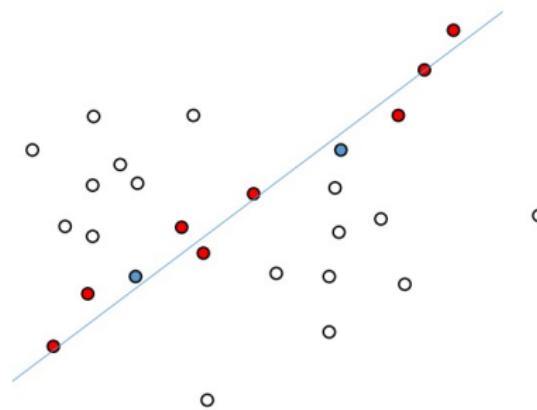
Line Fitting Using RANSAC

- Iteration 5



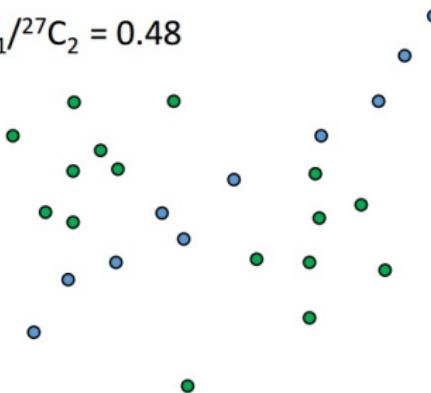
Line Fitting Using RANSAC

- Iteration 5



Why RANSAC Works?

- Inliers vs Outliers
- $P(\text{selecting outliers}) = {}^{17}C_2 / {}^{27}C_2 + {}^{17}C_1 {}^{10}C_1 / {}^{27}C_2 = 0.48$



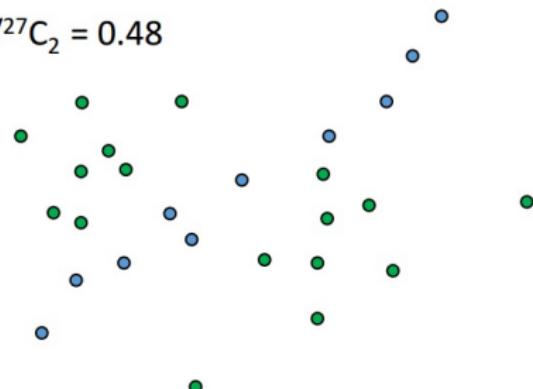
Why RANSAC Works?

- Inliers vs Outliers

- $P(\text{selecting outliers}) = {}^{17}C_2 / {}^{27}C_2 + {}^{17}C_1 {}^{10}C_1 / {}^{27}C_2 = 0.48$

- After 5 iterations...

- $P(\text{selecting outliers}) = (0.48)^5 = 0.026$



Why RANSAC Works?

- In general:
- $p = 1 - (1 - w^n)^k$

Where,

p = probability for selecting inliers

w = ratio of inliers to total #points

n = minimum #points required (for line = 2, circle = 3)

k = #iterations

RANSAC Algorithm

Determine:

n —the smallest number of points required (e.g., for lines, $n = 2$,
for circles, $n = 3$)

k —the number of iterations required

t —the threshold used to identify a point that fits well

d —the number of nearby points required
to assert a model fits well

Until k iterations have occurred

Draw a sample of n points from the data
uniformly and at random

Fit to that set of n points

For each data point outside the sample

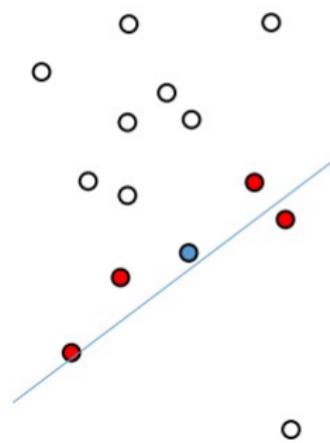
Test the distance from the point to the structure
against t ; if the distance from the point to the structure
is less than t , the point is close

end

If there are d or more points close to the structure
then there is a good fit. Refit the structure using all
these points. Add the result to a collection of good fits.

end

Use the best fit from this collection, using the
fitting error as a criterion



Algorithm 10.4: RANSAC: Fitting Structures Using Random Sample Consensus.

Slide Credits

Radhakrishna Dasari